

RWTH AACHEN UNIVERSITY
Chair of Computer Science 2
Software Modeling and Verification

Master Thesis

**Compilation of Quantum Programs with Control Flow
Primitives in Superposition**

Sascha Thiemann
Matr.-No.: 406187
Study Program: Computer Science M.Sc.
October 19, 2024

Supervisors: apl. Prof. Dr. Thomas Noll
Chair for Software Modeling and Verification
RWTH Aachen University

Prof. Dr. rer. nat. Dominique Unruh
Chair for Quantum Information Systems
RWTH Aachen University

Contents

| | | |
|----------|------------------------------------|-----------|
| 1 | Introduction | 1 |
| 2 | Background | 2 |
| 2.1 | Quantum Computing | 2 |
| 2.1.1 | Superposition | 3 |
| 2.1.2 | Entanglement | 4 |
| 2.1.3 | Quantum Gates | 4 |
| 2.1.4 | Measurement | 6 |
| 2.1.5 | Relevant Algorithms | 6 |
| 2.1.6 | Circuit optimization | 7 |
| 2.2 | Quantum Control Flow | 8 |
| 2.2.1 | Branching | 9 |
| 2.2.2 | Iteration | 10 |
| 2.2.3 | Limitations | 10 |
| 2.3 | Quantum Languages | 11 |
| 2.3.1 | Quantum Control Machine | 12 |
| 2.3.2 | OpenQASM Language | 14 |
| 2.4 | Compilation | 16 |
| 2.4.1 | Lexer (Lexical Analysis) | 16 |
| 2.4.2 | Parser (Syntax Analysis) | 17 |
| 2.4.3 | Semantic Analysis | 18 |
| 2.4.4 | Code Generation | 19 |
| 2.4.5 | Optimization | 20 |
| 2.4.6 | Tools | 21 |
| 3 | Concept | 23 |
| 3.1 | Language Overview | 23 |
| 3.1.1 | Blocks and Scopes | 23 |
| 3.1.2 | Data Types | 24 |
| 3.1.3 | Basic Operations | 25 |
| 3.1.4 | Measurements | 25 |
| 3.1.5 | Control Flow | 26 |
| 3.1.6 | Expressions | 26 |
| 3.1.7 | Composite Gates | 27 |
| 3.2 | Error Handling | 29 |
| 3.2.1 | Warnings | 29 |

| | | |
|----------|--|-----------|
| 3.2.2 | Critical Errors | 30 |
| 3.3 | Optimization | 31 |
| 3.3.1 | Optimization Rules | 31 |
| 3.3.2 | Circuit Graph | 33 |
| 3.4 | Command Line Interface | 36 |
| 4 | Implementation | 38 |
| 4.1 | Compiler | 38 |
| 4.1.1 | Symbols | 38 |
| 4.1.2 | Scope | 38 |
| 4.1.3 | Symbol Table | 38 |
| 4.2 | Lexical and syntactic analysis | 38 |
| 4.2.1 | Grammar | 39 |
| 4.2.2 | Data structures and classes | 40 |
| 4.3 | Semantic Analysis | 40 |
| 4.3.1 | Declaration Analysis | 41 |
| 4.3.2 | Type Checking | 41 |
| 4.3.3 | Error Handling | 41 |
| 4.4 | Code Generation | 42 |
| 4.4.1 | Expressions | 42 |
| 4.4.2 | Composite Gates | 42 |
| 4.5 | Optimization | 42 |
| 4.5.1 | Circuit Graph | 42 |
| 4.5.2 | Optimization Rules | 43 |
| 4.5.3 | Optimization Algorithm | 43 |
| 4.6 | Testing and Continuous Integration | 43 |
| 5 | Conclusion and Future Work | 44 |
| | References | 44 |

1 Introduction

2 Background

In the following section, we introduce and discuss different concepts that are referenced in later parts of this thesis. Firstly, we give a general introduction into quantum computing with some basic background knowledge on how quantum computers work and which quantum mechanical principles are essential for them. Additionally, we discuss more specific knowledge about quantum algorithms and optimization techniques for quantum circuits. Next, we discuss quantum control flow in more detail; this includes the formal definitions and its limitations. Then, we review existing quantum programming languages in general and some specific examples. Lastly, we give an overview on the topic of compilation and the different phases of a compiler.

2.1 Quantum Computing

While computers are prevalent and important in today's society, there are many relevant problems which classical computers cannot currently and perhaps will never realistically be able to solve. Quantum Computing (QC) is gaining more momentum as the technology that could solve at least some of these problems. For example, Quantum algorithms like Shor's algorithm [Shor97] could provide a significant improvement for prime factorization given sufficient technology. Therefore, it is estimated to be a valuable market with many of the largest technology companies as well as governments investing billions in the research and development of quantum technology [RDB*22, Pres18]. In the following section, we take a look at the basic concepts of a quantum computer and the core principles it relies on.

Classical Computers are based on simple operations executed on bits, like **and**, **or**, and **not**. These bits can either have a value of 0 or 1. Similarly, at their core, quantum computers apply simple operations, like **controlled not**, and **Hadamard**, on quantum bits (qubits). On a higher level, a classical computer executes operations on a register, consisting of multiple bits while a quantum computer operates on quantum registers, consisting of multiple qubits. In contrast to classical bits, quantum computers use the unique properties of quantum mechanics to enable qubits to have not just one value of either 0 or 1 but a combination of both. The phenomenon, where a particle or qubit exists in a combination of both states, is called *superposition*. Additionally, quantum computers also use the idea of *entanglement* to their advantage. Two qubits are entangled when the value of one is dependent on the value of the other. The combination of superposition and entanglement enables quantum computers to solve specific problems more efficiently than classical computers [RDB*22].

Models for Quantum Computers can be divided into three main categories, the *analog model*, the *measurement-based model*, and the *gate-based model*. The analog model

uses smooth operations to evolve a quantum system over time such that the resulting system encodes the desired result with high probability. It is not clear whether this model allows for universal quantum computation or quantum speedup [DiCh20b]. Instead of smoothly evolving a system, the measurement-based model starts with a fixed quantum state, the cluster-state. The computation is accomplished by measuring qubits of the system, possibly depending on the results of previous measurements. While there are different measurement-based models, one technique to apply gates is to leverage quantum teleportation, so called gate teleportation [Jozs05]. The result is a bit-string of the measurement results [DiCh20b, Niel06]. Lastly, the gate-based model uses a digitized, discrete set of qubits that are manipulated by a sequence of operations represented by quantum gates. The result is obtained by measuring the qubits at the end of the computation. Although digital quantum computation is more sensitive to noise than analog computations, the digitization can also be used for quantum error correction [DMN13] and to mitigate the increased noise [DiCh20b]. Furthermore, because qubits are actively manipulated and not passively evolved, digital quantum computers are more flexible than analog ones [RDB*22]. Therefore, the gate-based model is the most common model and this thesis will mainly focus on it.

Intermediate measurements are also supported sometimes.

Possible section on: no cloning/deleting [WoZu82, KuBr00]

2.1.1 Superposition

The first important property of quantum mechanics used by quantum computers is the idea of superposition. Qubits in superposition are often informally described as simultaneously having a value of 0 and 1 until their state is measured. However, a qubit in superposition is more formally a linear combination of its basis states. The basis states are the states where the qubit has a value of 0, written $|0\rangle = \begin{pmatrix} 1 \\ 0 \end{pmatrix}$, and 1, written $|1\rangle = \begin{pmatrix} 0 \\ 1 \end{pmatrix}$ [DiCh20a]. Furthermore, the state can be reduced to a simple vector. Therefore, a state ψ in superposition can be written as:

$$|\psi\rangle = \alpha |0\rangle + \beta |1\rangle = \alpha \begin{pmatrix} 1 \\ 0 \end{pmatrix} + \beta \begin{pmatrix} 0 \\ 1 \end{pmatrix} = \begin{pmatrix} \alpha \\ \beta \end{pmatrix}.$$

The factors α and β are the amplitudes of the basis states and are complex numbers. The factors must also satisfy the condition $|\alpha|^2 + |\beta|^2 = 1$. This is a result of the relation between the amplitudes and the probability to which basis state the state will collapse when measured, described in Sec. 2.1.4.

Beside $|0\rangle$ and $|1\rangle$, there exist more relevant short hands for quantum state. For example, $|+\rangle$ and $|-\rangle$ are states in uniform superposition, i.e. both basis state are equally likely, and often used when discussing quantum state und transformations. They are defined as follows:

$$|+\rangle = \frac{1}{\sqrt{2}}(|0\rangle + |1\rangle) = \frac{1}{\sqrt{2}} \begin{pmatrix} 1 \\ 1 \end{pmatrix} \quad \text{and} \quad |-\rangle = \frac{1}{\sqrt{2}}(|0\rangle - |1\rangle) = \frac{1}{\sqrt{2}} \begin{pmatrix} 1 \\ -1 \end{pmatrix}.$$

2.1.2 Entanglement

Another important quantum mechanical concept is entanglement. Simply said, two qubits are entangled when their values depend on each other. An example would be a quantum system where two qubits are in superposition and equally likely to collapse to either 0 or 1; whichever value one qubit collapses to when measured, the second one will also collapse to the same value. Additionally, changes to one of the qubits can also affect the other one. This happens independent of the locations of the two qubits [RDB*22, HHHH09].

A more formal definition for an entangled state uses the definition of a composite system. Two separate quantum systems can be represented as a single system with the tensor product of both systems. For example, the combined state $|\psi\rangle$ of the separate states $|0\rangle$ and $|1\rangle$ can be represented as:

$$|\psi\rangle = |0\rangle \otimes |1\rangle = |01\rangle = \begin{pmatrix} 0 \\ 1 \\ 0 \\ 0 \end{pmatrix} = \begin{pmatrix} 1 \\ 0 \end{pmatrix} \otimes \begin{pmatrix} 0 \\ 1 \end{pmatrix}.$$

When a quantum state cannot be expressed as a tensor product of two states, the state is entangled. The previous example is a case of a maximally entanglement Bell state [DiCh20a, MHH19], often denoted β_{00} , and can be expressed as the following:

$$\beta_{00} = \frac{1}{\sqrt{2}}(|00\rangle + |11\rangle) = \frac{1}{\sqrt{2}} \begin{pmatrix} 1 \\ 0 \\ 0 \\ 1 \end{pmatrix}.$$

The entanglement of states is used by leveraging the effect of the qubits on each other to collaborate to calculate the result. Although this can be simulated on classical computers, it cannot be achieved “natively” because all classical bits are independent of each other. Moreover, quantum algorithms not using entangled states can often be simulated efficiently on classical computers [MHH19]. Therefore, entanglement is at the core of quantum computing but it can also have unintended consequences one needs to be aware of when designing quantum algorithms.

To calculate specific functions or intermediate values, quantum algorithms may need to use additional qubits or registers whose state can, in turn, be entangled with the main data of the algorithm. If this entanglement is not resolved in time by, e.g., uncomputing the changes to the qubit or register, it can interfere with future calculations or measurements and cause the results to be invalid. This effect is called *disruptive entanglement* [YVC24].

Uncomputing as a concept was not introduced before

Cannot find literature besides [YVC24] which calls this effect disruptive entanglement, use anyway?

2.1.3 Quantum Gates

In gate-based quantum computers, the transformations applied to the quantum data are represented by *quantum gates*. Similar to quantum states, which can be represented

by linear combinations of basis states, or vectors, quantum gates can be formulated as linear transformations of these combinations, or a matrix. Because the result of such a transformation also needs to be a valid quantum state, the transformation needs to be norm-preserving, or *unitary* [DiCh20a]. The most relevant and often used unitary gates are depicted in Tab. 2.1

| | Gates | Matrix | Ket-notation |
|---------------------|--------------|--|--|
| Pauli gates | X | $\begin{pmatrix} 0 & 1 \\ 1 & 0 \end{pmatrix}$ | $ 0\rangle \mapsto 1\rangle$ $ 1\rangle \mapsto 0\rangle$ |
| | Y | $\begin{pmatrix} 0 & -i \\ i & 0 \end{pmatrix}$ | $ 0\rangle \mapsto i\rangle$ $ 1\rangle \mapsto - i\rangle$ |
| | Z | $\begin{pmatrix} 1 & 0 \\ 0 & -1 \end{pmatrix}$ | $ 0\rangle \mapsto 0\rangle$ $ 1\rangle \mapsto - 1\rangle$ |
| Hadamard gate | H | $\frac{1}{\sqrt{2}} \begin{pmatrix} 1 & 1 \\ 1 & -1 \end{pmatrix}$ | $ 0\rangle \mapsto +\rangle$ $ 1\rangle \mapsto -\rangle$ |
| Phase gate | $P(\lambda)$ | $\begin{pmatrix} 1 & 0 \\ 0 & e^{i\lambda} \end{pmatrix}$ | $ 0\rangle \mapsto 0\rangle$ $ 1\rangle \mapsto e^{i\lambda} \cdot 1\rangle$ |
| Controlled-NOT gate | CX | $\begin{pmatrix} 1 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 \\ 0 & 0 & 0 & 1 \\ 0 & 0 & 1 & 0 \end{pmatrix}$ | $ 00\rangle \mapsto 00\rangle$ $ 01\rangle \mapsto 01\rangle$ $ 10\rangle \mapsto 11\rangle$ $ 11\rangle \mapsto 10\rangle$ |
| Toffoli gate | CCX | $\begin{pmatrix} 1 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 1 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 1 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 1 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 1 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 1 \\ 0 & 0 & 0 & 0 & 0 & 0 & 1 & 0 \end{pmatrix}$ | $ 000\rangle \mapsto 000\rangle$ $ 001\rangle \mapsto 001\rangle$ $ 010\rangle \mapsto 010\rangle$ $ 011\rangle \mapsto 011\rangle$ $ 100\rangle \mapsto 100\rangle$ $ 101\rangle \mapsto 101\rangle$ $ 110\rangle \mapsto 111\rangle$ $ 111\rangle \mapsto 110\rangle$ |

Table 2.1: List of relevant quantum gates in matrix representation as as functions in ket-notation.

A matrix U is unitary if it has an inverse matrix which is equal to its conjugate transpose U^\dagger , i.e. the following must hold:

$$UU^\dagger = I.$$

Therefore, all transformations applied to quantum states in a gate-based quantum computer must be reversible by definition. This limitation does not apply to classical computers where non-reversible transformations, e.g. mapping an arbitrary bit to a specific value, are easily implementable.

To design a useful quantum computer or language, the set of gates should be *universal*. A set of gates is universal if any gate can be simulated by a combination of

the gates from the set with arbitrary accuracy [BrBr02]. An example for a universal set of gates is the combination of the Toffoli gate together with the Hadamard gate [DiCh20a].

Add paragraph on implicit measurement

2.1.4 Measurement

For quantum computer to be of any use, we need a way to read out information about its state. However, the information we can obtain from a quantum system is limited by the quantum measurement postulate. The postulate states that the only way, to gain any information from a quantum system, is to measure it. When measuring a quantum state, the state irreversibly collapses to one of its basis states. Furthermore, this is a probabilistic transformation and the original state in superposition cannot be recovered from the result. Therefore, in contrast to all other transformations, measurements are neither unitary nor reversible. For a state $|\psi\rangle = \alpha|0\rangle + \beta|1\rangle$, the measurement collapses the state to $|0\rangle$ with a probability of $|\alpha|^2$. Correspondingly, the state will collapse to $|1\rangle$ with a probability of $|\beta|^2$ when measured [DiCh20a].

Measurement can be represented as a measurement basis set $\{M_i\}_i$ which requires the following condition:

$$\sum_i M_i^\dagger M_i = I.$$

The probability that outcome i is obtained when measuring a state $|\psi\rangle$ is equivalent to $|M_i|\psi\rangle|^2$. After the measurement of outcome i , the state $|\psi'\rangle$ will be equivalent to

$$|\psi'\rangle = \frac{M_i|\psi\rangle}{|M_i|\psi\rangle|} = \frac{M_i|\psi\rangle}{\sqrt{\Pr[\text{observe } i]}}.$$

2.1.5 Relevant Algorithms

Since quantum computers differ greatly from classical computer not only in their technology but also in the concepts they use for calculation, they cannot function without specially designed algorithms. These algorithms need to exploit the special quantum properties of qubits to achieve *quantum advantage*, i.e. a better complexity than any classical algorithm. One of the first algorithms to show its quantum advantage was the Deutsch–Jozsa algorithm [DeJo92]. Deutsch et al. define a problem that can be solved in exponential time on classical computer and present a quantum algorithm which can solve the problem in polynomial time. The Bernstein-Vazirani algorithm [BeVa93] is another example with shown quantum advantage, resulting in a polynomial speed up. However, currently, there does not exist a use case for either of the algorithms and, therefore, they are only of limited theoretical interest [DiCh20c].

An algorithm with more potential for practical use is Shor’s algorithm [Shor97]. It presents a more efficient, polynomial-time quantum implementation for the discrete logarithm, i.e. find r for a given a, x, p such that $a^r = x \pmod p$. The algorithm is of special interest because Shor also provides a reduction of prime factorization to order finding; order finding is a special case of the discrete logarithm where $x = 1$. Modern

cryptography is often based on the complexity of factoring large prime numbers, e.g. the commonly used RSA cryptosystem [RSA78]. Therefore, an advanced quantum computer could break these systems with Shor’s algorithm [MVZJ18]. Not only does this prospect provide a practical use-case for QC but it also results in the research field of *post-quantum cryptography* [BeLa17].

Another relevant algorithm or transformation is the quantum Fourier transform (QFT) [Copp02]. Beside being used as a subroutine in Shor’s algorithm, it is also relevant for other algorithm, e.g. addition of quantum registers [Drap00]. Similar to the discrete Fourier transform [Wino78] which operates on vectors, the QFT_{2^n} operates on the quantum equivalent of vector, quantum registers, of size n . Registers of size n consist of n qubits. From the register, the QFT extracts the present periodic features. Then, other algorithms can use these features for their calculations.

2.1.6 Circuit optimization

Despite the expansive theoretical foundations for QC, the current state of the art for its technology is limited. However, the technology is nearing its first milestone towards useable quantum computers with the advent of prototypes with noisy intermediate-scale quantum (NISQ) technology [BFA22]. Nevertheless, the technology is still far away from fault-tolerant quantum computers and, by definition, limited in the number of available qubits. Furthermore, the gate count of NISQ era quantum computers is limited by the inherent noise which is increased with each additional transformation [Pres18]. Therefore, attributes such as the gate count of a quantum algorithm are an important metric for its utility. To improve the utility of an algorithm, its quantum circuit can be optimized with different techniques and rules.

There exist many kinds of optimization techniques for quantum circuits. They are mostly concerned with optimizing the gate count of quantum circuits with the use of peephole optimizations, as described in Sec. 2.4.5. These techniques can range from general rules [GaCh11, LBZ21], that can be applied to all quantum circuits, to hardware-specific optimizations [KMO*23]. Furthermore, machine learning based optimization frameworks for quantum circuits are also gaining popularity [FNML21, LPM*24, RLB*24].

The simplest general optimizations are so called *null gates* [GaCh11]. They are gate combinations or gates under specific conditions that are equivalent to the identity gate I . Therefore, any occurrence of such a null gate can be removed from the circuit. The most basic example for null gates is the double application of a self-inverse gate, i.e. a gate which is its own inverse. These include the H , X , Y , and Z gates. The identities are also depicted in Fig. 2.1. Furthermore, the same holds for any controlled version of a self-inverse gate such that the rule can also be applied to $CNOT$ and similar gates.

$$\boxed{H}\boxed{H} = \boxed{X}\boxed{X} = \boxed{Y}\boxed{Y} = \boxed{Z}\boxed{Z} = \boxed{I}$$

Figure 2.1: Null gates of self-inverse gates.

2 Background

The second kind of null gates are gates that do not have an effect under specific conditions. For example, a controlled gate U has no effect on its qubits if we know that the control is $|0\rangle$. Similarly, the X gate does not have an effect on a qubit in the $|+\rangle$ state. Hence, the two circuits depicted in Fig. 2.2 are null gates and semantically equivalent to an identity gate.



Figure 2.2: Null gates for gate in specific conditions.

Another class of optimizations are called *control reversal*. Control reversal describes gate combination equalities based on the symmetry of the controlled Z gate. For the controlled Z gate, it is semantically equivalent to apply the Z gate to the second wire with a control on the first and to apply the Z gate on the first wire with a control on the second one. Based on this and together with the equalities $HZH = X$, and $HXH = Z$, a controlled X gate surrounded by H gates on both wires can be represented as the reversed X gate. Both equalities are depicted in Fig. 2.3 and Fig. 2.4.




Figure 2.3: Control reversal of the controlled Z gate.




Figure 2.4: Control reversal of CX .

In contrast to general optimization rules, hardware-specific optimizations are mostly not concerned with the reduction of gates based on mathematically equal gate combinations; however, they exploit either the specific properties of the hardware for optimizations or replace gates with cheaper equivalents on the specific hardware. For example, a shuttling-based trapped-ion quantum computer operates by physically moving ions to segments in the hardware where operations can be applied. Since ions can and must be freely moved, swapping qubits can easily be accomplished by physically changing the position of their hardware equivalent. On the software side, this can be achieved by removing the swap-gate and swapping all following instances of both qubits. [KMO*23].

2.2 Quantum Control Flow

The idea of quantum control flow was first used by Altenkirch et al. [AlGr05] when defining a functional programming language with quantum control flow elements. The language uses an if-statement in superposition, if° , which is used to, e.g., defined the Hadamard gate as a function *had* instead of a matrix. The *had* function takes a qubit

as an input. If the qubit is true, i.e. the value is one, the function returns a uniform superposition of true and false, where true has a negative sign. Correspondingly, for a false input, a uniform superposition with both signs positive is returned.

$$\begin{aligned} \text{had} : Q &\rightarrow Q \\ \text{had} : x &\mapsto \text{if}^\circ x \\ &\quad \text{then } \{false \mid -true\} \\ &\quad \text{else } \{false \mid true\} \end{aligned}$$

Quantum control flow can be divided into *quantum branching* and *iteration* [YVC24]. In the following, we will discuss both branching and iteration in superposition as well as the limitations of quantum control flow.

2.2.1 Branching

Based on the work presented by Altenkirch et al. [AlGr05], the concept of quantum control flow, more specifically quantum branching, was expanded on and formally defined by Ying et al. [YYF12]. They introduce two different types of quantum branching, quantum guarded commands, and quantum choices as a special case of guarded commands. The definition of quantum guarded commands is based on Dijkstra's guarded commands [Dijk75]. Guarded commands concern the nondeterministic execution of functions based on Boolean expressions, where the nondeterminism derives from the possible overlapping of the guards. In contrast, quantum branching allows for execution of functions based on a value in superposition. The functions are executed such that the result may be a superposition of the results of the individual functions [YVC24]. Quantum branching is, e.g., used in simulation algorithms like [BGB*18] and [LoCh19]. Furthermore, many basic concepts such as controlled gates can be represented as quantum branching or even single qubit gates as seen in the previous example of the Hadamard implementation.

The formal definition for classical guarded commands is given by:

$$\Box_{i=1}^n b_i \rightarrow C_i$$

where C_i is a command guarded by a Boolean expression b_i . The command can only be executed if the expression is true. Similarly, quantum guarded commands map to a set of quantum programs P_i . Further, a set of qubits or quantum registers and a corresponding orthogonal basis $|i\rangle$ is given. However, the set of qubits guarding the program must be disjoint from the set of qubits used in the program. Without this condition, the resulting operation may not be unitary. For example, an X-gate that is executed if the wire it operates on is 1 always results in a values of 0; therefore, the operation can be neither reversible nor unitary. The resulting quantum guarded command is of the following form:

$$\Box_{i=1}^n \bar{q}, |i\rangle \rightarrow P_i.$$

The quantum programs are guarded by the basis states and the control flow results from the superposition of these basis states [YYF12].

2.2.2 Iteration

Quantum iteration can be implemented either as quantum recursion or quantum loops. While some languages implement loops based on the measurement of qubits or registers [Ying11], the concept of quantum iteration requires the body of the loop to be executed in superposition based on a guard in superposition [YYF12].

While classical iteration takes an operation and repeats it on a classical register for k iterations, quantum iteration is dependent on a value k' in superposition and, correspondingly, returns a quantum register in superposition. Moreover, it is a special case of quantum branching and heavily restricted by the limitations of quantum computers [YVC24].

2.2.3 Limitations

While quantum control flow is often based on the corresponding control flow primitives on classical computers, it is restricted by multiple limitations imposed by quantum computers. Therefore, many control flow primitives that are used in classical programs can either be not used at all or in a limited capacity. There are two main limitations for quantum programs. Firstly, all gate-based quantum computers need to adhere to *reversibility*. Secondly, programs need to follow the *synchronization* principle for them to return any useful results [YVC24].

Reversibility

As introduced in Sec. 2.1.3, any sequence of instructions on gate-based quantum computers, excluding measurements, is required to be reversible by definition as they are all unitary transformations. Therefore, any quantum control flow is also required to adhere to this principle. A resulting limitation, that is not present on classical computers, is that any guards for guarded commands need to be immutable in the commands themselves. For example, if a qubit's state is flipped when its value is 0, the resulting command will always return value of 1. When a program returns the same result regardless of which statements were executed, the program cannot be reversible. This limitation is also inherent in the definition of quantum guarded commands, as described in Sec. 2.2.1. Moreover, control flow, as implemented in classical computers, is also not possible. At the most basic software level, modern computers use jump and conditional jump instructions to implement branching and loops. However, any classical jump instruction is inherently irreversible. Not only can a jump go to a section of code that is accessible without any jumps, multiple jumps can also lead to the same line of code. Therefore, a reversed program cannot know which path was taken in original program [YVC24].

A simple solution seems to be offered by the *Landauer Embedding* [Land61]. Fundamentally, the idea of the embedding is to turn a non-reversible function into a reversible one by not only returning the output but also the input of the function. For example, for a domain D and a codomain D' , any non-reversible function $f : D \rightarrow D'$ can be given as a reversible function $g : D \rightarrow D' \times D$ with $g(x) = (f(x), x)$. In the case

a quantum program with, e.g., jump instructions based on the Landauer embedding, the output would be the result of the program and a complete history of which path was taken through the program. However, because the quantum data depends on the program history, they become entangled. This leads to disruptive entanglement, as described in Sec. 2.1.2, causing invalid results [YVC24].

Synchronization

As previously discussed, reversibility alone is not the only limiting factor on quantum control flow. When handling control flow, similar to the classical implementation, with a program counter in superposition, the program counter can become entangled with the data and result in disruptive entanglement leading to an invalid result. To avoid this issue, the program must not only be reversible but also adhere to the principle of *synchronization*. It states that control flow must become independent from the data. Further, because any quantum program needs to be synchronized to return any useful results, while loops dependent on a value in superposition need to be bounded by a classical value [YVC24].

2.3 Quantum Languages

With the emergence of quantum computing, many quantum languages were introduced. Most languages focus on a lower level representation of quantum circuits. An example is the popular Open Quantum Assembly Language (OpenQASM) [CBSG17]. OpenQASM consists mainly of quantum and classical registers that can be manipulated by predefined and composite gates. Additionally, some classical control flow is possible with if-statements depending on classical bits or measurements. As its name suggests, the language is designed for low level interactions with quantum computers and mostly used to directly describe a quantum circuit. In Sec. 2.3.2, OpenQASM is discussed in more detail.

In contrast to the low level circuit descriptions of OpenQASM, there are also languages with a focus on high level interactions. One such language is Tower [YuCa22]. It does not only allow for basic qubits and registers in superposition but also abstract data structures such as lists. Another example is the language Silq [BBGV20] which allows for the automatic and safe uncomputation of registers after they have been used for, e.g., intermediate calculations. What both languages have in common is the restriction to quantum data while using only classical control flow.

Although quantum control flow was formally defined by Ying et al. [YYF12], as described in Sec. 2.2, over ten years ago, only very few languages have incorporated the principle. One example is the functional programming language proposed by Altenkirch et al. [AlGr05] where quantum branching is used to define, e.g., the Hadamard gate. Only recently was the Quantum Control Machine with quantum control flow at its core proposed by Yuan et al. [YVC24]. It presents an instruction set similar to classical assembly languages but for quantum computers and discusses the resulting

limitations for the language. In the following section, we will discuss the quantum control machine in more detail.

2.3.1 Quantum Control Machine

The Quantum Control Machine (QCM), proposed by Yuan et al. [YVC24], is an instruction set architecture that does not only allow for data in superposition but also quantum control flow. The architecture is designed around the limitations of control flow in superposition.

The syntax and logic of the QCM are both heavily influenced by classical assembly languages. Similar to classical computers, the language provides a finite set of quantum registers which are all initialized to a value of zero. The instruction set of the architecture does not only provide limited gate transformations and swap operations but also more classical operations on registers such as get-bit operations and simple arithmetical operations like addition and multiplication. However, what makes the QCM stand out are the jump instructions that enabled quantum control flow.

The gates of the architecture are limited to the X and Hadamard gate H . However, since the QCM enables quantum branching, any gate can become a controlled gate such that the X gate can easily be used in combination with quantum branching to create a Toffoli gate. Together with the Hadamard gate, the gate set is therefore universal, as described in Sec. 2.1.3.

There are three kinds of jump instructions. The first is a simple jump based on a given offset, the second is a conditional jump that performs a basic jump when a given register is 0, and, lastly, an indirect jump which is based on the value of a given register. Although the jump instructions are based on jumps in classical computers, they are limited by the restriction of unitary gates and must adhere to *reversibility* and *synchronization* [YVC24], as described in Sec. 2.2. An overview of some QCM instructions is depicted in Tab. 2.2.

| Operation | Syntax | Semantics ¹ |
|-------------------|------------------------|---|
| No-op | nop | Only increases instruction pointer by the branch control register. |
| Addition | add $ra\ rb$ | Adds register rb to ra . |
| Multiplication | mul $ra\ rb$ | Multiplies register ra by rb . |
| Jump | jmp p | Increases branch control register by p . |
| Conditional Jumps | jz $p\ ra$ | Increases branch control register by p if ra is 0. |
| | jne $p\ ra\ rb$ | Increases branch control register by p if ra is not equal to rb . |

Table 2.2: An excerpt of the QCM instruction set with instructions used in later examples.

¹After all operations, the instruction pointer is increased by the value of the branch control register.

When quantum computers are based on unitary gates, all their operations need to be unitary and, therefore, reversible as well. This limits quantum jump instructions and prohibits them to work like their classical equivalent. However, the problem of a reversible architecture and instruction set is not unique to quantum computers but was also taken into consideration for classical architecture to, e.g., increase energy efficiency of classical computers [AGY07, TAG12]. To enable reversible jumps, the QCM adapts the *branch control register* from the reversible Bob architecture [TAG12]. Instead of directly changing the instruction pointer of the machine, the branch control register specifies how much the instruction pointer advances after each instruction.

The branch control register can then be manipulated reversibly by, e.g., adding or subtracting from it. To jump by a given *distance*, the branch control register needs to be increased to *distance*. However, after the instruction pointer has reached the desired location, the register needs to be decreased to its original value. Otherwise, the pointer would continue to jump in larger increments and any further jumps, i.e. modification to the branch control register, would not jump to the correct location. Since the jump instructions are defined to be reversible, the instruction set also includes a reverse jump instruction which instead decreases the branch control register by a given offset. Therefore, a jump instruction always requires a reverse jump instruction to reset the program counter. Similarly, other operations can also be represented as the reverse operation of an existing one. For example, subtraction can be implemented as reverse addition. Further, to make the code easier to read and write, the QCM also allows for named labels, which can be used for jump instructions instead of offsets. The offset to the given label can then be computed at compile time.

An example of a classical program and the reversible equivalent can be seen in Fig. 2.5 and Fig. 2.6 respectively. Both programs calculate x^y for two registers x and y . While the first example has classical jumps that are not reversible, the second example uses reversible jump instructions and their reverse counterpart to create a reversible algorithm.

```

1      add    res $1
2      add    r1 y
3  l1:  jz     l2  r1
4      mul    res x
5      radd   r1 $1
6      jmp    l1
7  l2:  nop

```

Figure 2.5: A non-reversible exponentiation algorithm.

```

1      add    res $1
2      add    r1 y
3  l1:  rjne   r11 r1 y
4  r12: jz     l2  r1
5      mul    res x
6      radd   r1 $1
7  r11: jmp    l1
8  l2:  rjmp   l2

```

Figure 2.6: Reversible exponentiation algorithm.

Although such a program counter addresses the issue of reversibility, it can become entangled with data registers when in superposition. This can lead to disruptive entanglement where the output of the program becomes invalid [YVC24]. To prevent any disruptive entanglement of the data and control registers, QCM programs must

2 Background

adhere to the principle of synchronization, as described in Sec. 2.2. It requires that the control flow is separated from the data at the end of execution. However, this is not the case for the reversible example program in Fig. 2.6 which, therefore, is not a valid QCM program.

The issue, that occurs in the loop of the reversible example, is the *tortoise and hare* problem. Given a superposition of two different values a and b in the y register, the loop will execute a and b times respectively. Therefore, one of the two loops will finish before the other. Since we must adhere to synchronization, the instruction pointer needs to become independent of the two values again. However, because the branch with the faster execution of the loop cannot simply wait, the other branch cannot catch up and the instruction pointer cannot become independent of the data values. Consequently, the program does not adhere to synchronization. To prevent this issue, the program must include padding operations which are executed instead of the main loop. Furthermore, the loop also needs to be bounded by a classical value, as described in Sec. 2.2.3. The results in an algorithm, as depicted in Fig. 2.7, that calculates $x^{\min(y, \max)}$. Here, \max is a classical bound to the number of loop iterations, as required.

| | | | | |
|----|-------------------|-------------------|------------------|------------------------|
| 1 | <code>add</code> | <code>res</code> | <code>\$1</code> | |
| 2 | <code>add</code> | <code>r1</code> | <code>max</code> | |
| 3 | <code>l1:</code> | <code>rjne</code> | <code>r11</code> | <code>r1 max</code> |
| 4 | <code>r12:</code> | <code>jz</code> | <code>l2</code> | <code>r1</code> |
| 5 | <code>r13:</code> | <code>jg</code> | <code>l3</code> | <code>r1 y</code> |
| 6 | <code>mul</code> | <code>res</code> | <code>x</code> | |
| 7 | <code>r14:</code> | <code>jmp</code> | <code>l4</code> | |
| 8 | <code>l3:</code> | <code>rjmp</code> | <code>r13</code> | |
| 9 | <code>nop</code> | | | <code>; padding</code> |
| 10 | <code>l4:</code> | <code>rjle</code> | <code>r14</code> | <code>r1 y</code> |
| 11 | | <code>radd</code> | <code>r1</code> | <code>\$1</code> |
| 12 | <code>r11:</code> | <code>jmp</code> | <code>l1</code> | |
| 13 | <code>l2:</code> | <code>rjmp</code> | <code>r12</code> | |

Figure 2.7: A synchronized, reversible exponentiation algorithm.

2.3.2 OpenQASM Language

The Open Quantum Assembly Language (OpenQASM) 3 [CJA*22] is the successor of the OpenQASM 2 [CB SG17] language. Both languages are imperative and machine independent quantum languages. They are low level quantum languages and, thereby, concretely describe a quantum algorithm in the form of a circuit. OpenQASM 2 developed into a de facto standard and is often used as an intermediate language for different quantum tools [CJA*22]. OpenQASM 3 was developed to fit the changing needs of current quantum research and hardware while being mostly backwards compatible except for some uncommon cases. For example, some keywords were added or changed for the successor such that identifiers of OpenQASM 2 circuits may be invalid

in the successor language. Since OpenQASM 3 is the new and improved standard, we will focus on its features in the following section.

OpenQASM 3 requires the header to indicate the language in the circuit header for any top-level circuit. This is achieved by adding `"OpenQASM 3.0";` to the beginning. Additionally, the language supports the inclusion of other source files which can be included with the `include` keyword.

Similar to other quantum languages, OpenQASM operates on two basic data types. The first is the classical bit while the second is the qubit. Both primitives can also be used in registers with a fixed size. Additionally, OpenQASM 3 also supports further classical data types such as angles and signed and unsigned integers. In contrast to its predecessor where any identifiers have to start with a lowercase letter, in OpenQASM 3, identifiers can start with a range of unicode characters with some exception.

The basic operations of the language can be divided into unitary and non-unitary operations. In OpenQASM 3, all unitary operations are based on the unitary $U(a, b, c)$ where a, b, c are angular parameters. While OpenQASM 2 supported a controlled-NOT gate natively, the successor requires the gate to be defined with, e.g., the NOT gate and a control modifier. The control modifier can be used to turn any arbitrary unitary gate into a controlled gate with an arbitrary number of control qubits. Therefore, the formerly predefined gate CX must now be defined by the programmer or represented by a NOT gate with a control modifier, e.g. `ctrl @ x`. Lastly, the non-unitary operations are `measure` and `reset`. While the `measure` operation measures the state of a qubit and saves it to a classical bit, the `reset` operation discards the value of a qubit and replaces it with the $|0\rangle$ state.

The programmer can not only use the operations and modifiers provided by OpenQASM 3 but can also define custom gates. These user-defined gates are defined with an identifier for the gate and a fixed number of single qubit arguments and angular parameters. In the body of the gate definition, the user can apply a sequence of gates to the qubit arguments with the given angular parameters. Additionally, the language also provides implicit iteration. This means that the application of a single qubit gate to a quantum register will be interpreted as separate applications of the gate to all qubits in the register.

In Fig. 2.8, an example circuit, written in OpenQASM 3, is depicted. The circuit takes two qubits, brings them into an entangled superposition, measures their state and saves the result to a classical register. In the beginning of the circuit definition, the circuit header indicates the language and the X , CX , and H gate are defined based on the predefined unitary U . Then, the quantum and classical register, both with a size of 2, are defined. Next, the Hadamard gate H is applied to the first qubit in the quantum register followed with the application of a controlled-NOT gate to both qubits. Lastly, the state of both qubits is measured and the result is saved to the classical bits.

```

1  "OpenQASM 3.0";      /* Indicate language in circuit header. */
2
3  gate x a { U(pi,0,pi) a; }          /* Define x gate. */
4  gate cx a, b { ctrl @ x a, b; }    /* Define cx gate. */
5  gate h a { U(pi/2, 0, pi) a; }     /* Define h gate. */
6
7  qubit[2] reg;           /* Definition of quantum register. */
8  bit[2] res;            /* Definition of resical register. */
9
10 h reg[0];              /* Apply h gate to fist qubit in register. */
11 cx reg[0], reg[1];     /* Apply cx gate to the qubits. */
12
13 res[0] = measure reg[0]; /* Measure qubit and save to bit. */
14 res[1] = measure reg[1]; /* Measure qubit and save to bit. */

```

Figure 2.8: Code for an OpenQASM 3 example circuit.

2.4 Compilation

The execution on a computer is controlled by a program. This program is written in a specific language unique to the hardware of the computer, machine code. However, this language is often neither human readable nor suitable for writing complex systems. Therefore, most programs are written in a more accessible language. The program can then be translated to the machine code with a *compiler*.

A compiler translates a program written in a source language to a program in a target language. The compilation process can be divided into multiple steps. The first step is the *lexical analysis* to transform the source code into a sequence of tokens. Next, the syntactic structure of the code is analyzed by the *parser*. Then, the code is *semantically analyzed* to find semantic errors and infer information for the following phases. Lastly, the *code generation* step generates the code in the target language. Additionally, the compiler may perform optimizations on the code before generating the target code or it may *optimize* the resulting target code [Oliv07, VSSD07]. In the following, we discuss the different steps of a compiler individually.

2.4.1 Lexer (Lexical Analysis)

The lexical analysis of the source program takes the character stream and groups together associated characters producing a sequence of tokens [Oliv07]. Therefore, the step is also referred to as *tokenization* [Gref99]. The process can be divided into the *scanning* and *screening* of the character and token sequence [DeRe74].

The scanning process groups together substrings into textual elements, or tokens. In contrast to the characters and substrings, these tokens have defined meanings and may have additional attributes. For example, they may include identifiers, operator, comments, and spaces. In the case of the identifier token, an additional attribute could be the string value of the identifier. They can be specified with the help of a regular

grammar or regular expression [DeRe74, VSSD07].

[VSSD07] is extensive book, cite specific chapter somehow?

After being divided into a sequence of tokens, the screening step drops any characters or sequences of characters not relevant to the compilation from the program code. These may include characters such as spaces and tabs, or white space in general, and character sequences such as comments. Further, it may also recognize additional special symbols, such as keywords, and map them to a designated token. For example, a identifier with a value of “while” could be mapped to the corresponding token of the `while`-token.[DeRe74].

Some example regular expressions for a lexical analysis are depicted in Fig. 2.9. The code depicts regular expressions for integers, identifiers, comments, and white space in ANTLR syntax. The integer can either be an arbitrary sequence of characters between zero and nine without a leading zero or just zero with a length of at least one. Similarly, an identifier is a sequence of lower and upper case alphabetical characters, numbers, and underscores with a length of at least 1 and without a leading number. In contrast, a comment is any string starting with a double slash until the line break and white space is any white space characters. Additionally, the comment and white space also define a scanning step where both are discarded.

add reference to section discussing ANTLR

```

1 INTEGER      : [1-9] [0-9]* | '0' ;
2
3 IDENTIFIER    : [a-zA-Z_] [a-zA-Z_0-9]*;
4
5 COMMENT       : '//' ~[\r\n]* -> skip;
6
7 SPACE         : [ \t\r\n\u000C] -> skip;
```

Figure 2.9: An example of a regular grammar for the lexical analysis.

2.4.2 Parser (Syntax Analysis)

The lexical analysis of the compiler yields a sequence of tokens with a known meaning; the structure of the program, however, is not apparent in the token sequence. For example, an operator-token does not indicate what the operands are. To gain knowledge of the structure of the program, the parser step analyzes the syntactic structure of the source program and creates a parse tree from it. The compiler can then use the tree by, e.g., walking over it to generate the target code. This step should also detect and report any syntactical errors, like a missing closing parentheses [VSSD07].

While the lexical analysis can be achieved with regular expressions, the syntactic structure of a program must be represented by, at least, a context-free grammar. Since regular languages are a subset of context-free languages, the parsing step can also perform the lexical analysis. However, there are multiple reasons why the lexical and syntax analysis are separated. Firstly, the separation of both analysis makes the compiler more modular and extensible. Furthermore, using regular expression for the

citation needed?
(Chomsky, "Three models for the description of language")

2 Background

lexical analysis prevents it from being more complex than necessary with a context-free grammar. Lastly, the lexer can be more efficient when generated from regular expressions instead of a context-free grammar [VSSD07].

There exists two main kinds of parsing a grammar, either top-down or bottom-up. Top-down parsing creates a parse tree based on an input sequence of tokens starting from the root and creating the nodes in a depth-first approach. It yields a left-most derivation for the input sequence and can be implemented as a recursive-descent parser. The most common form of top-down parsing is *LL*-parsing, where the input is read from left to right, yielding a leftmost derivation. To improve the efficiency of parsers, the context-free grammar is often restricted such that it can be parsed without backtracking with a fixed length *lookahead* onto the token sequence. Such grammar are called *LL(k)* grammars where *k* is the length of the lookahead [VSSD07, PaFi11].

In contrast, bottom-up parsing builds the parse tree from the leaves up to the root. Furthermore, instead of yielding a left-most derivation, it produces a right-most derivation. Similar to top-down parsing, the most common bottom-up parsers scan the input from left to right which are therefore LR-parsers. Moreover, they can also be implemented more efficiently when restricting the grammar to a maximum lookahead. These grammar the *LR(k)* grammars [VSSD07, PaFi11].

An example grammar for parsing simple integer expressions is depicted in Fig. 2.10. Similar to the regular expressions in Fig. 2.9, the grammar is given in ANTLR syntax. An expression is either the sum of another expression and a term or just a term. In turn, a term is either the product of a term and a factor or just a factor. Lastly, a factor is either an expression in parenthesis or an integer. Here, the definition of an integer is omitted. However, it can be seen in the previous example. The grammar is defined such that a generated parse tree inherently adheres to the order of operations.

```
1 exp      : exp '+' term      | term;
2
3 term     : term '*' factor   | factor;
4
5 factor   : '(' exp ')'       | INTEGER;
```

Figure 2.10: An example of a context-free grammar for parsing simple expressions.

2.4.3 Semantic Analysis

The parser analyzes the syntactic structure of a program with a context-free grammar; however, an analysis without any context is not sufficient for an analysis of non-syntactic, i.e. semantic, constraints of the program. This step is performed by the semantic analysis. The semantic analysis is used to throw semantic errors that may prevent the program from being compiled such as the use of undefined identifiers. Further, it may also enforce constraints that prevent runtime error such as type checking in a strongly typed language. Additionally, the analysis step may also process and

save declarations and similar information to a symbol table which can be used in the code generation or optimization [Oliv07, SWW*88]. Moreover, the semantic analysis may not only throw errors but can also be used to infer additional information for further compilation steps. For example, besides preventing operations on operands with invalid types, the analysis may deduce which operation to apply to the operands based on their type; in the case of two integers, the analysis may infer an integer additions for the “+”-operator while two floating point values require floating point operations [Wait74, VSSD07].

What specifically the semantic analysis does is dependent on the design of the language being analyze. For example, a loosely typed language may have limited type checking, when compared to a strongly typed language, if any at all. Further, the implementation of the analysis can differ greatly. However, all implementations have some common elements. It requires the propagation of attributes through the syntactic structure of the program to enable the analysis. In the case of type checking, the analysis must pass on the type of a variable. Moreover, it does not only need to know the types of variables and constants, i.e. leafs in a parse tree, but also the resulting type of an expression using them. For example, a integer added to a floating point value may result in a floating point value. To infer and propagate these attributes, the parse tree may need to be transverse [Wait74, VSSD07].

2.4.4 Code Generation

After the semantic analysis of the program, which, at this stage, is in the form of a parse tree, the compiler can generate the code. Here, the compiler can either generate the target code, e.g. machine code, directly or translate the parse tree into an intermediate code. The translation of the source code to the intermediate can be thought of as the *frontend* of the compiler, with the translation of the intermediate to the target being the *backend*. While the intermediate code will need to be translated again into the target language, the use of an intermediate representation can increase the modularity and extensibility of a compiler. Additionally, it can also ease the construction of a new compiler. When creating a new compiler from a source language to a target language the front end of an existing compiler for the source can be combined with an existing compiler to the target if both are using the same intermediate language [VSSD07, GFH82].

em dash (—) here?

The most common issues when generation the target code are the evaluation order of expressions, register and storage allocation as well as related issues, context switches, and instruction selection [GFH82]. While these issues are critical for compilers that translate classical languages, i.e. not quantum languages, to machine code, they are mostly not relevant for the translation of quantum computers, since quantum computers do not offer same features and abstractions that classical computer do; they have, e.g., no storage, other than the quantum registers. Therefore, we will not discuss these issues in more detail.

citation needed?

2.4.5 Optimization

While the lexical, syntax, and semantic analysis combined with the code generation are the essential parts of a compiler, without which it would not work, the optimization step is also important. It used to apply either machine-independent or machine-dependent optimizations. The optimizations can be applied to the parse tree, a possible intermediate representation, and the generated target code depending on the optimization itself. While the removal of unreachable code, e.g. code after a return statement, can most easily be performed on the parse tree, machine-dependent optimizations can, more appropriately, be performed on the target or intermediate code [Oliv07, VSSD07].

Two collaborating machine-independent optimizations that are often applied by compilers are *constant propagation* and *constant folding*. Constant propagation analyzes the code to find variables with constant values throughout all executions and replaces the variables in, e.g., expressions with their corresponding constant value. By itself, constant propagation may only result in marginal improvement, loading a constant literal instead of the values of a variable; however, in combination with constant folding it can result significant improvement. Constant folding evaluates expressions or subexpressions with constant values at compile time, resulting in less calculations at runtime. This can significantly increase the performance of a program especially if large expressions or expressions in loops can be folded. Propagating constant values through the code enables more constant folding and, therefore, can improve its effectiveness [WeZa91].

Another optimization technique, that can work in tandem with constant folding and propagation, is *loop unrolling*. When executing a loop, each iteration needs to check the halting condition and possibly execute an increment statement which can result in significant overhead. Furthermore, since the condition is checked before each iteration, the different executions cannot be executed in parallel. To prevent or reduce the performance overhead from these issues, the loop body can be executed multiple times and the increment statement adjusted accordingly. Further, if the number of iterations is constant, the loop can be removed entirely and replaced by the repeating loop body [HuLe99]. In this case, constant propagation and unrolling can help evaluate the halting condition such that the loop can be unrolled.

Similar to loop unrolling, *function inlining* also replaces some part of the code with an equivalent code body to reduce overhead. The inlining of a function replaces a function call with the function body. This mitigates the overhead caused by the function call. Additionally, the inlining also enables or simplifies further optimization such as constant folding and propagation [TGS22]. However, excessive function inlining may significantly increase the code size and, therefore, have negative effects on the caching of the code and, in turn, on the performance. This effect can be decreased or mitigated by other transformations or code reductions enabled by the inlining [PeMa02].

Lastly, *peephole optimization* is concerned with optimizing inefficient code patterns by analyzing mostly the machine or intermediate code and removing or replacing them. The pattern can be replaced with a reduced number of instructions that have the same effect or instructions that are easier to execute. For example, loading a constant value

of 0 and executing an addition has no effect and can be removed while a multiplication by 8 can be replaced with a cheaper left shift by 3. Furthermore, depending on the language operated on, peephole optimization can also implement some rudimentary constant folding; loading two constants A, B and adding them can be replaced by loading the constant $A + B$. These patterns can be saved in tables and systematically applied to the program code [McKe65, TvS82].

2.4.6 Tools

A compiler is a complex program that can not only require a lot of coding but also is also likely to include errors when written from scratch. While the earlier syntactic stages use general algorithms, writing a custom lexer and parser can require a substantial workload and is prone to errors. Therefore, there exist many tools that can either help create a lexer and parser or generate them entirely [PaFi11, ZLY17]. In the following, we can briefly discuss different available compiler generation tools.

Two tool for compiler generation, often used in tandem, are the Fast Lexical Analyzer (*Flex*)² and GNU *Bison*³ [DoSt99]. The Flex is a lexer generator while Bison is a general-purpose parser generator. Both generators target C and C++ code with Bison having the experimental feature to support Java. Both tools are an extension and improvement of previous tool, *Lex* and *Yacc* respectively. Bison implements a bottom-up parser and can parse most grammars. However, the tool is optimized for $LR(1)$ grammars, i.e. bottom-up parsing with a lookahead of a single token [ZLY17, Aaby03, DoSt99].

While Flex and Bison are separate tools for the lexing and parsing of program code, *ANTLR*⁴ [PaQu95] combines both purposes into a single tool. ANTLR stands for “ANother Tool for Language Recognition” and can be used to generate lexers as well as parsers. In contrast to Bison, it implements top-down parsing and can recognize any $LL(k)$ grammar with $k > 1$. Additionally, while Flex and Bison are mainly targeting C and C++, ANTLR can generate lexers and scanners in a variety of languages, including C++, Java, Python, and C#. With its newest version ANTLR4, the lexing and parsing rules can be given in the form of a context-free grammar with the terminals of the grammar given as regular expressions. At the beginning of each grammar, the name of the grammar is given and it is indicated whether the grammar describes a lexer, parser, or a general grammar, possibly containing both. The grammar rules can be either lexer or parser rules. They always begin with a rule name, followed by a colon and the different alternatives, separated by a vertical bar, and terminated with a semicolon. While parsing rule names are given in lowercase letters, lexer rule names must begin with upper case letters.

A simple ANTLR4 grammar is depicted in Fig. 2.11. Firstly, it indicates that this grammar describes a general lexer-parser combination and gives it the name “simple_exp”. Then, a parsing rule called “expression” is defined where the expression

add information about *op* =, what does it do, why is it used

²<https://github.com/westes/flex>

³<https://github.com/akimd/bison>

⁴<https://github.com/antlr/antlr4>

2 Background

```
1 grammar simple_exp;  
2  
3 expression : INTEGER op=OPERATOR expression  
4             | INTEGER;  
5  
6 OPERATOR   : '+'  
7             | '-';  
8 INTEGER    : [1-9] [0-9]* | '0' ;
```

Figure 2.11: Simple ANTLR4 grammar for expressions.

is either just an integer or an integer, and operator and another expression. Lastly, the operator and integer lexing rules are defined. The operator is either the “+” or “-” character and the integer is a regular expression for an arbitrary sequence of number characters without a leading 0.

indicate that *op* = is used

3 Concept

In this chapter, we give a high-level overview of the different concepts and feature provided by our proposed programming language Luie and the corresponding compiler. Furthermore, we discuss some of the design decision and draw comparisons to other quantum languages.

We start of with an overview of the language. This includes the overall structure, the different data types and statements, and some special behavior. Next, we discuss the error handling of the compiler. There, we introduce and explain all possible warnings and critical errors that can error. The error handling is followed by the discussion of the optimizations that the compiler may apply to the code; this includes some inherent optimizations as well as optional peephole optimizations. Lastly, we discuss the command line interface of the compiler, its options and their corresponding behavior.

3.1 Language Overview

In the following section, we discuss the different concepts that the language provides for writing quantum programs; we discuss their behavior and some use cases for them. Additionally, we explain why some features of other languages were not incorporated and the reasons for some special behaviors. Firstly, we discuss code blocks and their corresponding variables contexts, or scopes. Then, we introduce the data types that are available in the language. Next, the basic operations as well as the more complex control flow statements are listed and explained. This is followed by an overview of the possible expressions and the available operations and build in functions. Lastly, we discuss the composite gates of the language.

The following description is a high level overview of the concepts and functionalities of the language. The concrete syntax and implementation of these features is discussed in Ch. 4.

3.1.1 Blocks and Scopes

Similar to many other languages, Luie uses code blocks and corresponding scopes. Both the code blocks and scopes are used to structure the code and enable the reuse of identifiers in different contexts. Each Luie program is contained in the main code block. This main code block can contain arbitrarily many other code blocks. In turn, these code blocks can also contain any number of nested blocks. However, the main code block can only exist once and is the parent block of all others. Therefore, it also cannot depend on any other block. The main block not only differs from the others in terms of hierarchy but it is also the only block that can contain composite gate

3 Concept

definitions. The composite gates are defined at the top of the main code block and are followed by other declarations or statements.

Similar to the main code block, all other code blocks also consist of declarations and statements. A new code block is defined either in the definition of a composite gate or in the body of a control flow primitive; this includes the if-statement, else-statement, and for-loop. Each code block has a unique scope.

Scopes represent the variable context of a code block. Since each scope corresponds to one code block, they are also hierarchically structured. In turn, the parent of a scope is the scope that directly contains this scope, while ancestors are all scopes that also indirectly contain the scope. In contrast, the descendant of a scope are all scopes that are directly and indirectly contained in the scope. Furthermore two scopes are independent if one is neither an ancestor or descendant of the other. In a given scope, the program can access all variables previously defined in this scope and all its ancestors, however not the variables of any descendants. In contrast, two independent scope do not have access to the variables defined in the other scope. Therefore, two independent if-statements can define a variable with the same name in their scope. Additionally, a scope can also overwrite the definition of a variable of an ancestor with its own definition while the ancestor's definition is not affected.

bad formulation

same language can be applied to code blocks

3.1.2 Data Types

Luie is mainly focused on being a quantum language with quantum control flow; this is also reflected in its data types. The language operates mostly on two types of quantum data, registers and qubits. Both behave as described in Sec. 2.1 where a register represents an array of qubits with a fixed length. Both registers and qubits are declared similarly. First, the `qubit` keyword is used to indicate the declaration. In the case of a register, the size of the register is given in brackets. Lastly, the identifier for the register or qubit is specified. For classical data types in most language, the variables need to be explicitly initialized. In quantum computers, however, qubits can only be initialized to $|0\rangle$. Therefore, Luie does not require or allow for an initialization value for qubits and registers but always defaults to $|0\rangle$ for all qubits and register entries.

Here or in implementation: while to the outside register array of qubits, internally qubit special case of register (with length 1).

The second class of data types in Luie are classical data types. They can be divided into constant values, constant variables, and loop iterators. In contrast to most classical language with similar data type, they are not referenced in the compiled program; they are only used at compile time and, therefore, must all have a constant value of known behavior. Constant values are the most simple data types and represent any value given directly in the program code. Currently, they can only be integers.

expand on this idea

Similar to constant values, constant variables, often just called constants, also have a constant values that is calculated when they are declared. However, they can not only be an integer but also unsigned integers and floating point values. Furthermore, they can be given by an expression when declared and reused as often as needed in the context they were declared in. Constants are declared with the `const` keyword and need to be given an identifier as well as the type of the constant. The type is either the

`int`, `uint`, or `double` keyword. Lastly, an expression for the value of the constant is given.

The loop iterator data type is declared when creating a loop and is also only accessible in the loop body. Its value corresponds to the current loop iteration, e.g. for the first iteration 0 and the last one $n - 1$ where n is the number of iterations. Loop iterators are usually not implemented as a special data type but by iterating, e.g., an integer; however, Luie does not allow of mutable, classical data type and, therefore, cannot increment an integer to iterate over a loop. In turn, it requires a special data type which value is constant for each iteration.

3.1.3 Basic Operations

Any language needs some basic operations to manipulate its data type. While these operations may be addition, multiplication, and similar ones, quantum computers manipulate their data with gates. Luie provides multiple predefined gates with which the qubits and registers can be manipulated. The syntax for applying gates to qubits is similar to other quantum language; the name of the gate is given and all parameters are listed next, separated with commas.

gate-based quantum computer, not generally

The first kind of gate set natively available are the basic single qubit gates; these include the X , Y , and Z gates as well as the Hadamard gate H . The logical extension of single qubit gates is the second gate set, the multi-qubit gates; they include the controlled-not gate CX and the Toffoli gate CCX . These gates are not required for the overall gate set to be universal because they can easily be simulated with the control flow abilities provided by Luie. However, they are included because most programmers are used to them. Finally, the parameterized gates are the last gate set. In contrast to the basic single and multi-qubit gates, they are not constant in their behavior but depend on a parameter. This parameter is given as an expression and is evaluated at compile time; it is given in parenthesis after the gate itself. Currently, the only parameterized gate available is the phase gate $p(\lambda)$. The respective behavior and mathematical definition of all predefined gates in Luie is discussed in Sec. 2.1.3. Since the language provides both the Toffoli and Hadamard gates, the overall gate set is universal. Furthermore, any gates that are not predefined, can easily be implemented as composite gate; they are described in Sec. 3.1.7.

bad reason, write better one!

However, one feature not provided for gate applications is implicit iteration. Implicit iteration allows gates that operate on a single qubit to be applied to an entire register instead by giving the register as a parameter. Since Luie provides an loop statements with which a gate can explicitly be applied to entire register with additional control over the application, implicit iteration is not implemented in the language.

3.1.4 Measurements

- Short paragraph about how measurements are handled/added

3.1.5 Control Flow

The main focus of the proposed language is quantum control flow. Therefore, Luie provides two different control flow statements, a loop statement with a fixed number of iterations and an if statement operating on data in superposition.

The loop statement can be used to iterate over a code block a fixed number of times. While the number of iterations needs to be known at compile time, it can depend on constants known at compile time such as the size of a register. However, it can, e.g., not depend on the measurement of a qubit or even the value of a qubit in superposition. This relates back to the principle of synchronization, as described in Sec. 2.2.3. Additionally, the body of the loop can depend on iteration index, i.e. the current value of the iterator. Therefore, the loop statement can, e.g., iterate over a register or part of it and apply a series of gates to each qubit. Furthermore, the loop body can also include all statements and declarations available for each code block. In turn, the loop statement can contain nested loops or parts of its body can be guarded by an if statement. However, the target language of the compiler and quantum computers in general do not have a concept of local variable. Therefore, while declarations can be performed, a declaration in a loop body is translated to a “global” declaration of a qubit for each iteration. This can significantly increase the qubit count of the program. To optimize the current implementation, a single qubit could be declared that is reset or measured at the end of each iteration and, thereby, reducing the required number of qubits.

In contrast to loop statements, the if statements do not depend on constant data known at compile time; their execution, however, depends on the value of a qubit in superposition. Further, the guard of the statement can only be a single qubit, not a register, and cannot include boolean expressions. The body of the statement can contain any other statements, including further if and loops statements, as well as declaration. When compiled, each gate application inside the body is guarded by the given qubit and translated to a controlled version of the gate where the guard is the control. However, the declaration is not dependent on the value of the guard. Therefore, similar to declarations inside loop statements, the declaration is always performed independent of the guard’s value. In turn, the declaration of qubits or register inside an if statement can increase the qubit count but not to the same extent as the declaration in loop statements because it is only performed once and not for each iteration. Additionally, the if statement also offers an optional else case besides the main code block. In this case, the translated statements of the else block are controlled by the negated guard.

explain why no boolean expressions?

3.1.6 Expressions

The language allows for complex expressions that are evaluated at compile time. These expressions can be used to access specific indices of a register or define the range of an iterator used by a loop statement. Besides the typical operations like addition and multiplication, the language also implements different functions. However, the

targeted quantum computers should not need to and, in some cases, cannot evaluate these expressions, their values need to be not only constant but also evaluated at compiler time; they can neither depend on the measurement nor the value of a qubit.

The first kind of operations are the basic operations; they include addition and subtraction, multiplication and division, and the negation. The grammar of the language is designed such that the operations inherently follow the order of operations. Additionally, parenthesis can be used to adjust this order. The operations operate either on an integer value or identifier. If a integer is given, it can easily be parsed to an integer and used with the operations. However, the identifier can either be a constant or an iterator. For either, the compiler looks up the current value in the symbol table and executes the other operations.

Besides basic operations, the language provides more complex functions that can be used for specific calculations. Firstly, the `sizeof` function takes a register as an input and returns the size of the given register; this can be useful for creating a loop statement that iterates over the size of a register. Furthermore, it can be used to create registers with the same size of another. Next, the `power` function takes two expressions as parameters, evaluates them and raises the first parameter's value to the value of the second parameter. Lastly, similar to the `power` function, the `min` and `max` functions both take two expressions as parameters and return the minimum and maximum value of the evaluated expressions respectively.

Luie provides two different expression for creating ranges. Since they create ranges, or more specifically loop iterators, and the other operations only operate on basic numeric types, e.g. integers and floating point values, they are not grouped together with the others but implemented separately. The first way to define a range is to give the start and end values separated by two dots, e.g. `1..9` for iterating from 1 to 9. In this case, both the start and end values are included in the iteration. However, this expression can only take integers for the start and end values not expressions; this limitation is imposed such that the expression remains readable. The other operation for creating ranges is the `range` keyword followed by the start and end values in parenthesis and separated by a comma, e.g. `range(1,9)`. For this expression, the parameters can also be given as expressions. Furthermore, the `range` expression also provides a short hand, where only the length `len` of the range is given. In this case, it will always start from 0 and iterate to `len - 1`. This case is most handy for creating ranges to iterate over registers.

Lastly, in contrast to many other languages, the access of a register is not implemented as an expression. A register access, i.e. a qubit, is not a classical data type and cannot be used in any context other than as a gate parameter or as an if statement guard. Therefore, it is implemented as a separate grammar rule and not as an expression.

limitations: no implicit casting, currently all values in expression same type

3.1.7 Composite Gates

Similar to OpenQASM, Luie also allows for the declaration of composite gates. They can be used to declare a custom gate that applies the gate combination specified in

3 Concept

the gate body. This can be particularly useful for gate combinations often used in the code to reduce redundancy and, further, improve the readability of the code. For example, the swap gate can be defined on two qubits where three controlled-not gates are applied such that the values of both qubits are swapped. In this case, not only are three gates reduced to only one but the code clearly indicated that the values are swapped without prior knowledge on the effect of the controlled-not gates.

While the concept of a composite gate is also available in OpenQASM, Luie expands on the possibilities provided by OpenQASM's implementation. Firstly, the gate does not only allow for simple gate application statements but also for the control flow statements provided by Luie; this includes the loop statements and if statements in superposition, only qubit and register declarations are prohibited in the body of gate declarations. In contrast to composite gates in OpenQasm, not only qubits are allowed as parameter but also registers. Therefore, registers can be given as parameters and the loop statement can iterate over them resulting in gates that can depend on the length of a register and do not need to be reimplemented for each different size. However, in contrast to OpenQASM, the language currently does not allow of parameterized composite gates.

```
1 // Swaps the values of two qubits
2 gate swap(a, b) do
3     cx a, b;
4     cx b, a;
5     cx a, b;
6 end
7
8 // Performs a discrete Fourier transform on a register of qubits
9 gate qft(reg) do
10     for i in range(sizeof(reg)) do
11         h reg[i];
12         for j in range(sizeof(reg) - (i + 1)) do
13             qif reg[j + (i + 1)] do
14                 p(1/(power(2, (j + 1)))) reg[i];
15             end
16         end
17     end
18     for j in range(sizeof(reg) / 2) do
19         swap reg[j], reg[sizeof(reg) - (j + 1)];
20     end
21 end
```

Figure 3.1: Luie gate definition for the Quantum Fourier Transform.

Two examples for composite gates are depicted in Fig. 3.1. The first gate is a simple definition for the commonly used, and often predefined, `swap` gate. Below, a gate definition for the quantum Fourier transform is given. It not only references the previous swap gate, but also uses both loop and if statements. Firstly, the gate iterates over all qubits in the given register `reg` with size n . In each iteration, the gate

first applies the Hadamard gate to the qubits and, again, iterates over all remaining qubits in the register and applies the phase gate depending on the value of the qubit in superposition. Additionally, the parameter for the phase gate is calculated based on the offset of both indices. After iterating through the entire register and applying the corresponding gate, the gate swaps each qubit at index i with the qubit at index $n - i$ for $0 \leq i \leq \lfloor \frac{n}{2} \rfloor$.

3.2 Error Handling

Generally, an important part of a program is error handling; useful and precise error messages are essential for comfortable interactions with the program. This is especially the case for compilers where the user should not only easily understand what the issue is but also where in the source code the error occurred.

Our compiler has two types of errors with different severity. The first type is the warning. A warning from the compiler can indicate issues in the source code that may cause unintended behavior. However, the issues itself does not prevent the compilation of the program and is simply an indication that there may be something wrong. In contrast, the critical error is caused by a flaw in the source program that prevents the correct compilation and will result in the abortion of the compilation. In the following, we will discuss the different warnings and critical errors, the compiler may raise and how they are caused. Furthermore, we discuss why they are either a warning or critical error.

3.2.1 Warnings

The compiler can throw two different kinds of warnings. The first is the invalid range warning and the second is the unused symbol warning.

An invalid range warning can occur in the context of for loops. They iterate over a range that is defined by the user. It can be given as either a size n and iterate from 0 to $n - 1$ or a start and end index, i_{Start} and i_{End} respectively, and iterate from the start to the end. However, the range iterator is designed to only increase. Therefore a range where $i_{Start} \geq i_{End}$ is invalid. Since the for loop is unrolled at compile time, a range with a size less than or equal to zero can just be ignored. This however may be unintended behavior. Therefore, the compiler warns the user that the range is invalid.

The unused symbol warning is raised when a symbol, e.g. a register of composite gate, is defined in the source code but never used. The unused symbol does not have any negative effect on the compilation and the optimization step can easily remove, e.g., an unused register. Therefore, this is only a warning and the program can be compiled. However, an unused symbol may indicate, that the wrong symbol was used somewhere else or part of the program is no longer used, hence, warning the user of the unused symbol may prevent unintended behavior.

3.2.2 Critical Errors

The list of critical errors, often just referred to as errors, is more extensive than the list of warnings. It includes the invalid access, number of arguments for gates and functions, and size errors. Furthermore, there are errors for the attempted declaration of a variable that is already declared, an undeclared and type error, and, lastly, an error for the invalid use of a qubits in a guarded code block.

The first error is the invalid access error. It occurs when a register is accessed an index i , i.e. i is either smaller than 0 or larger than $size - 1$, where $size$ is the size of the register. While this error could easily be ignored and would cause no issue when compiling the program, the resulting code would be an invalid circuit description.

Secondly, the invalid number of arguments error for gates and functions is caused when the number of arguments given to either a gate or function is not as expected. For example, an the Hadamard gate, when not controlled, always expects one argument while the controlled-not gate requires two. Similarly, the `sizeof`-function operates on exactly one argument. The compiler cannot proceed when given too few arguments; for the opposite case, dropping any leftover arguments would result in an unexpected behavior. Additionally, too many arguments indicate an issue in the program. Therefore, the compiler reports an error for both cases and aborts the compilation.

Another error is the invalid size error. It occurs when a register is declared with an invalid size. A size is invalid if it is less or equal to 0. A register with non entries cannot be used for anything and, likely, indicates an issue in the program while a register with a negative amount of entries is impossible. Therefore the compilation is aborted and the error is thrown.

The next two errors are concerned with the declaration of variables in a given context; they are the undeclared and already-declared errors. An undeclared error is raised when a variable is used in a context where it is not defined. In this case, the symbol table does not have a symbol stored for the given identifier and the compiler cannot continue. In contrast, the already-declared error occurs when a variable is declared in a context where the same variable identifier has already been assigned to a different symbol. While the compiler could overwrite the previous declaration, this can easily lead to unexpected behavior and, in turn, most languages do not allow a declaration in the same scope to be overwritten.

The type error is thrown when a variable is used in a function or gate but is not the required type. Languages with loose typing may be able to convert some types to the required type by, for example, parsing the integer value of a string. However, this can not only result in unexpected behavior and hard to debug errors in the code but, in the case of a quantum language, it also would require the conversion from classical to quantum data which is not easily achievable.

Finally, the last error is the use-of-guard error; it occurs when a qubit is referenced in the context that is guarded by itself. While the compiler can easily translate any such occurrence, they result in an invalid circuit description. As described in Sec. 2.2.1, a gate that operates on and is controlled by the same qubit cannot be reversible. Therefore, the compiler prevents the generation of an invalid circuit and aborts the

change, dont like formulation

compilation.

3.3 Optimization

Our compiler does not only translate the source language Luie to the target language OpenQASM but can also apply optimizations to the program. Since quantum programs must be deterministic and, therefore, most language constructs are translated at compile time, many classical optimizations are inherently applied. The following, inherent optimizations are discussed in more detail in Sec. 2.4.5.

The first kind of inherent optimizations are constant propagation and constant folding. While the target language does allow for expressions, they need to be constant. Furthermore, some language features that depend on expressions have no equivalent in the target language so that they, and their corresponding expressions, need to be evaluated at compile time; this is the case for, e.g., the control statements. Therefore, any variable is always constant and its value is propagated to be used in the evaluation of expressions. Further, each expression is evaluated at compile time such that constant propagation and constant folding are inherently applied.

Secondly, loop unrolling is always applied to all loop statements. OpenQASM does not allow for loop statements or any other method for the iteration of commands, besides the restrictive implicit iteration. Furthermore, quantum computer in general cannot provide the ability to iterate over gates since they operate on static circuits. Therefore, to allow for loop statements in our language, the loop is unrolled entirely at compile time. Additionally, for each iteration the current values of the loop iterator is propagated as a constant through the loop body.

The last inherent optimization is function inlining. Our language provides the ability to define custom gates that consist of an arbitrary combination of gate applications and, possibly, control flow statements. While OpenQASM has a similar functionality, the composite gate OpenQASM provides are more restrictive; for example, they do not allow for registers as arguments to the gate. Furthermore, in contrast to classical computers, quantum computers do not natively support function calls or similar. Therefore, each time a composite gate is used, the compiler inline the gate body at the call site of the gate.

3.3.1 Optimization Rules

After the code is translated, the compiler can perform additional peephole optimizations; these are not inherent to the translation of the program and, therefore, are optional. The peephole optimizations can be applied to the internal representation of the source language and most hardware-independent optimizations are usually applied to the intermediate representation instead of the target language. However, the peephole optimization rules only operate on sequences of gate applications such that control flow statements or composite gates may only hinder, not aid, the effectiveness of the optimization. The overall performance of the optimization may be increased by optimizing composite gates before inlining their code so that a gate that is called ten

3 Concept

times only needs to be optimized once; nevertheless, inlining the gates before applying the rules may enable more optimizations and, thereby, increase the effectiveness at the cost of performance. Therefore, we apply the optimization rules after the code is translated. Additionally, the user can not only specify whether optimizations are applied but also which to apply when using the compiler. In the following section, we discuss the different optimization rules that can be applied. Additionally, the theoretical background for the rules is discussed in Sec. 2.1.6.

Why do we need to apply to translated code and not intermediate? most optimizations typically on intermediate

The first kind of optimization rules are the null gate optimizations; they describe sequences of gate applications such that the resulting behavior is equivalent to applying the identity gate. In the case of classical computers, an example is the sequential execution of two negation operations. In contrast, a quantum example is the application of two successive Hadamard gates. While these optimizations can easily be performed by the programmer themselves for a simple list of instructions, the manual optimizations increase in complexity when using composite gates and control flow statements. Moreover, the programmer cannot remove two null gates that are contained in two different successive composite gates by hand. Therefore, the optimization rules cannot only help to reduce the workload of the programmer but apply optimization rules are cannot be implemented without major changes to the program.

Next, the “peeping control” optimization rule also belongs to the null gate rules. However, its implementation requires some additional evaluations and, therefore, we separate it from the other null gate rules. The peeping control rule can remove a controlled gate from the circuit if value of the control wire is $|0\rangle$ at the position of the gate. To estimate the value of the control wire, the implementation needs to iterate over all previous gates on the wire. Therefore, while it is still a null gate, its implementation differs greatly from the other null gates. Furthermore, we implement an additional optimization that removes the control from the gate if the value of the control wire is known to be $|1\rangle$. While the rule does not reduce the gate count of the circuit, it can enable further optimization of the control wire. For example, two X gate on the control wire can be separated by a controlled-not gate. If the value of the control is $|1\rangle$ when the controlled-not gate is applied, the control can be removed and the now successive X gates can be removed with a null gate optimization. Since this optimization does not remove the gate from the circuit, it is not a null gate optimization.

Another optimization rule that the compiler implements is the “Hadamard reduction” rule. It implements the matrix equivalences for both the X and Z gate being surrounded by Hadamard gates, $HXH = Z$ and $HZH = X$. Thereby, the rule reduces the gate count of the circuit if the optimization rule is applied. However, the optimization rule does not remove the gate combination but replaces it with another; therefore, the rule is not a null gate optimization. Similar to the null gate optimizations, the rule itself is not hard to apply by hand. However, in combination with the more complex statements available in the language, the application of the rule is not trivial and the compiler can optimize parts of the circuit that would have required major changes to optimize them manually.

Dont like the description, cannot come up with better right now

Lastly, the control reversal optimization rule optimized a controlled-not gate that is

surrounded by Hadamard gates on all wires. When applied, the four Hadamard gates are removed and the control and target qubits of the controlled-not gate are switched. Therefore, the gate count is reduced by four gates. As described in Sec. 2.1.6, the optimization is based on two Hadamard reductions and the control reversal of the controlled Z gate. However, the reversal of the controlled Z gate does not have any direct gain when applied; it can only enable other optimizations. Furthermore, the application of the controlled Z reversal may also disable other optimizations. Therefore, an optimization algorithm using the rule would need to either test both possibilities or estimate the value of the application and, in turn, increasing the complexity of the optimization algorithm significantly. Because of this, our optimizations do not include the controlled Z reversal but implement only the special use case of the rule, the control reversal optimization rule.

3.3.2 Circuit Graph

While it is possible to directly apply optimizations to the program code or internal representation of that code, this approach can be tedious and error prone. For example, the easiest approach would be to iterate over the code and search for code sequences with more efficient but equivalent alternatives, similar to the peephole optimization pattern on classical computers presented in Sec. 2.4.5. However, two consecutive gates operating on a single wire may be separated by multiple gate applications on different wires in the programmatic description. Therefore, many simple optimization rules may not be applied when using a simplistic algorithm. A more complex approach would be to subdivide the program into lists of gate application where the wires being operated on by each list are disjoint. While this approach can result in the application of more optimization rules, it will also miss possible applications and already requires a complex implementation. Furthermore, improving on this method only increases its complexity and, in turn, makes it more prone to errors. Therefore, the language does not directly apply the optimizations to the program but uses a circuit graph description, based on the graph described by Kreppel et al. [KMO*23], to apply the optimizations.

The circuit graph is a graphical description of a quantum graph; it is an acyclic and directed graph. For each qubit in the circuit, there exist both an input node and an output node. Furthermore, each input node has exactly one outgoing vertex while each output node has exactly one incoming vertex. Besides the input and output nodes, all other nodes represent gates in the circuit. All gate nodes the number of incoming vertices is equivalent to the number of outgoing vertices. Additionally, the number of incoming, or outgoing, vertices is the same as the number of parameters of the gate. Importantly, in this case, any qubits controlling the application of the gate, e.g. the first qubit in a controlled-not gate, do also count to the number of parameters. In turn, there exists a vertex pair of an incoming and outgoing vertex each that represent a qubit to which the gate is applied. Therefore, for each qubit, there exists one path from its input node to its output node such that all gates that are applied to it are visited in order of application.

Graph Construction

When using the circuit graph to optimize the quantum program, the first step is to systematically construct the graph from the program. The creation of the graph starts with the input and output nodes for each qubit in the circuit. For each declaration of a qubit in the program, a input and output node pair is created. If a register is declared, a pair is created separately for each qubit in the register, i.e. for a register with size n , n pairs are created in total. Next, the gate applications in the program can be iterated. For each gate application, a corresponding gate node is created. To insert this node into the graph, each qubit parameter requires an incoming vertex to this node and a corresponding outgoing vertex from the node. Additionally, the incoming node must come from the gate that was previously applied to the qubit or the input node if no gate was applied yet. Similarly, the outgoing gate must lead to the either the next applied gate or the corresponding output. Therefore, for each qubit, the vertex coming into the output node can be diverted to the gate node and an outgoing vertex from the gate to the output node can be created. Repeating this step for all gate applications in the program result in the circuit graph.

An example of a simple, unoptimized circuit graph is depicted in Fig. 3.2. For simplicity, all applied gates are depicted in their corresponding node of the circuit graph. The circuit consist of three qubits, q_0 , q_1 , and q_2 ; their input and output nodes are depicted on the left and right of the graph respectively and are labeled with the corresponding qubit. Firstly, an X is applied to the first qubit q_0 . Than, a controlled-not gate is applied to the first two qubits q_0 and q_1 , where q_0 is the controlling qubit. Simultaneously, two Hadamard gates H are applied to the third qubit q_2 . Lastly, another X gate is applied to the first qubit q_0 while another controlled-not gate is applied to the second and third qubits q_1 and q_2 . In this case, the second qubit q_1 is the controlling qubit.




Figure 3.2: An example of a simple, unoptimized circuit graph.

Graph Optimization

The next step in the optimization process is the application of optimization rules, in this case peephole optimizations, to the circuit graph. To accomplish this, subgraphs of the graph need to be iterated systematically. Then, each subgraph is checked for an equivalent but more efficient or cost-effective alternative. If an alternative is found, the subgraph is replaced with it. In our case, we iterate over all possible subpaths of all qubit paths up to a certain length; this length depends on the maximum number of nodes that are affected by an optimization rule. A qubit path is a path that starts at the input node of a qubit and ends in the corresponding output node. Furthermore, it visits only the nodes corresponding to the gates that are applied to the qubit in the correct order. However, optimizing parts of the circuit may enable further optimizations. For example, after removing a gate combination, two previously separated Hadamard gates may now represent a null gate combination which can, in turn, also be removed. Therefore, the process needs to be repeated until no more optimizations were applied.

An example for this is the optimization process of the circuit graph depicted in Fig. 3.2. Firstly, subpaths for the qubit path of the first qubit q_0 are iterated. Here, the subpath of node it the path with the first two gate nodes X and CX . Since the X node the child of the input node, we know that the value of the qubit will be $|1\rangle$ after the application. Therefore, we know that the controlled-not gate will always apply the X gate to the second qubit q_1 and we can replace the CX gate node with a simple X node that is only visited by the second qubit's path. For simplicity, we assume that the optimizations are only applied after all subpaths are iterated such that there are no optimization that can be applied to the second qubit in this round of the optimizations. However, in our later implementation, the optimizations are applied while iterating over the circuit. Lastly, on the third qubit wire, there are two consecutive Hadamard gates that can be removed. Overall, the first round of optimizations results in the circuit graph depicted in Fig. 3.3.




Figure 3.3: Circuit graph after the first optimization.




Figure 3.4: Completely optimized graph.

In the next optimization round, the two X gates on the first qubit wire are now

3 Concept

applied consecutively. Therefore, they can both be removed from the circuit. Finally, the combination of an X gate child of an input node followed by a controlled-not gate can, again, be optimized such that the CX gate is replaced with a simple X gate on the third qubit wire. The result is a circuit where the first qubit remains unchanged and only an X gate is applied to the second and third qubit. This result is also depicted in Fig 3.4.

Graph Translation

After all possible optimizations rules were applied and no others were enabled in turn, the only remaining step is to translate the circuit graph back to a programmatic description. The qubits for the circuit are easily declared by iterating over all input or output node. Additionally, the qubit count can be reduced by leaving out all qubits where the output node is the child of the input node after the optimization. In the case of the optimized circuit, depicted in Fig. 3.4, the first qubit q_0 can be skipped when adding the declaration to the translated program.

For the gate applications, the only requirement is that, if a gate node is a descendent of another, its translated gate application statement must come after the statement of the ancestor node. Therefore, for most circuit graph there are multiple possible programs describing the circuit correctly. For example, when translating the completely optimized graph described above, the order of the X gate applications to both the second and third qubit is irrelevant such that both possibilities, i.e. either applying the gate to q_1 or q_2 first, are valid translations of the program. Furthermore, while iterating over all input or output nodes and adding the corresponding qubit and register declarations to the beginning of the program is the easiest approach, the order and placement of the declarations is also arbitrary as long as they are declared before their first use.

3.4 Command Line Interface

The command line interface of the compiler enables the programmer to interact with the compiler and specify specific behavior. There are currently four different parameters that can be specified, the input file, output file, optimizations to be applied, and the verbosity of the compilation.

The first, and most important, parameter is the input file. Since without a specified input file the compiler has no program to compile, the parameter is mandatory. It can be given by either `-i` or `-input` followed by the path to the input file. Similarly, the output file is specified with another parameter; it can be specified with `-o` or `-output` followed by the path to the output file. However, this parameter is not required and the default behavior is to create a file with the name `output.qasm` in the current directory.

Next, the optimization parameter can be used to specify which optimizations are applied to the code. Here, the possibilities are either none, the null gate, peeping control, Hadamard reduction, and control reversal optimization rules. Each optimization

has a specific keyword that can be passed as a parameter to indicate the optimization. Additionally, multiple different optimizations can be applied by listing them separated by a plus sign. For example, to apply both the null gate and peeping control optimizations, the correct parameter would be “`nullgate+peepingcontrol`”. The optimization parameter is specified by either passing “`-O`”, in this case an uppercase “o” to differentiate it from the output, or “`-optimization`”. When the optimization is not explicitly specified, the default behavior is no optimization.

currently, spaces are NOT allowed (should be implemented for ease of use)

Another parameter that can be passed to the compiler is the verbose parameter; it toggles the verbose mode of the compiler on. When this mode is active, the compiler prints info, warning, and error messages out to the user. In comparison, only error messages are displayed to the programmer when the compiler is executed without the verbose mode. In contrast to the other parameters, the verbose mode does not require any other arguments and can be activated by passing either “`-v`” or “`-verbose`” to the compiler. If the parameter is not passed, the compiler is executed normally.

Besides the parameters that can be passed to the compiler to modify its behavior, it also provides a help test. The help text can be accessed by passing the help parameter to the compiler with either “`-h`” or “`-help`”. Additionally, it is displayed when an invalid or unknown parameter was given by the programmer together with a message that indicates that the given parameter is invalid. The help text itself gives a short description of the compiler and lists all parameters that can be passed. Furthermore, each parameter is listed with a short description of its behavior.

An example for the usage of the CLI is depicted in Fig. 3.5. The first parameter is the input parameter; it indicates that the compiler should compile the Luie program called “`program.luie`” in the current directory. Next, the output file is specified. After the compilation is complete, the compiled code should be located in the build directory of the current directory and be called “`program.qasm`”. Following the specification for the input and output files, the optimizations parameter indicates which optimizations are to be applied to the program. In this case, the compiler applies both the null gate and peeping control optimization gates. Lastly, since the verbose parameter was not specified, the compiler will only print errors and no other informational messages or warnings.

```
1 ./LUIECompiler --input "./program.luie"
2                  --output "./build/program.qasm"
3                  --optimization nullgate+peepingcontrol
```

Figure 3.5: A command line interface example.

4 Implementation

- General introduction and information about the implementation
- Implementation using C# and .Net
 - Why use C# as language to implement compiler in?
 - High-level, easier, faster (coding not execution) and saver implementation than C, C++
 - More performant than Python or Javascript
 - More familiar with C# in general

4.1 Compiler

- Static compiler class
 - General information
- How are different parts called
- How do they interact?
- Discuss general concepts and classes used by compiler

4.1.1 Symbols

- List different kinds of symbols and there structure

4.1.2 Scope

4.1.3 Symbol Table

4.2 Lexical and syntactic analysis

- Describe grammar structure
- Different elements of grammar

The lexical and syntactic analysis steps of the compiler are not implemented by hand but ANTLR4 is used to generate both the lexer and parser. In Sec. 2.4.6, a shot introduction is given into the generation tool. The tool does not only generate the lexer and parser for the language but also some data structures, like the parse tree,

and abstract classes that are used for the following compiler steps. In the following, we discuss the ANTLR grammar of Luie and give an overview of the different classes that are generated by ANTLR.

4.2.1 Grammar

The entry point for the grammar is a parsing rule called **parse**. It states that the program consists of a main code block followed by the end of file lexer token **EOF**. The main code block starts with arbitrarily many gate declarations, or none at all, and is followed by any number of declarations and statements in any order, or none at all. Similarly, a “normal” code block also consists of any number of declarations and statements; however, it does not allow for gate declaration. The grammar for the entry point, main and general code block is depicted in Fig. 4.1

```

1 parse      : mainblock EOF;
2
3 mainblock  : gateDeclaration* (declaration | statement)*;
4
5 block      : (declaration | statement)*;
```

Figure 4.1: The basic structure parsing rules of Luie.

While syntactically it would be equal to say the main block consists of gate declarations and a code block, the code generation implementation requires the main code block to directly hold the declarations and statements. Internally, each code block object holds a list of declarations and statements which, in turn, can hold references to other codeblocks. However, a block itself cannot hold a reference to another block. Furthermore, the code generation is only called on the main code block which, then, recursively calls the code generation for each declaration and statement. Therefore, the main code block needs to contain the declarations and statements directly, otherwise the code generation will always generate an empty program.

There are three kinds of declaration rules in the grammar. The first is the declaration rule for qubits and registers. It starts with the register keyword lexer token, followed by an optional expression in brackets giving the size. If no size is given, a qubit is created; otherwise, a register with the given size is created. The identifier for the qubit or register is the last part of the rule.

Next, the constant declaration rule is used to declare constants that can be used at compile time, e.g. for the size of registers or to give the number of loop iterations. Similar to register and qubit declarations rule, the rule starts with a constant keyword token. Then, the identifier is given and separated from the type of the constant with a colon. Lastly the expression for the value of the constant is given after an equality sign.

The final declaration rule is used for composite gate. Again, it starts with a corresponding keyword and is followed by the identifier of the gate. Next, the optional gate

Should gate parameters be optional?

4 Implementation

parameter rule is given in parenthesis. The gate parameter rule matches any number of identifiers, separated by commas, requiring at least one identifier. Lastly, the gate declaration rule ends with a code block rule.

All declaration rules are given in Fig.4.2. Note that the depicted grammar is semantically equal to the grammar used in the implementation; however, some of the formatting and naming was adjusted for the thesis.

should this be specifically stated?

```
1 declaration      : registerDeclaration | constDeclaration;
2
3 registerDeclaration : REGISTERKEYWORD
4                   ('[' size=expression ']')?
5                   id=IDENTIFIER ';' ;
6
7 constDeclaration  : CONSTANTKEYWORD id=IDENTIFIER
8                   ':' type=TYPE '='
9                   exp=expression ';' ;
10
11 gateDeclaration   : GATEKEYWORD id=IDENTIFIER
12                   '(' param=gateParameter? ')'
13                   DO block END;
14
15 gateParameter     : IDENTIFIER (',' IDENTIFIER)*;
```

Figure 4.2: The parsing rules for declarations in Luie.

statements...

4.2.2 Data structures and classes

4.3 Semantic Analysis

The second step in our compilation process is the semantic analysis of the input program. The analysis checks for any semantic errors in the code that can be detected before the next compilation step, the code generation. Our implementation of the semantic analysis differentiates between a declaration analysis and type checking. The declaration analysis is concerned with checking for errors related to the declaration of variables. For example, when a variable is used, it check whether it was declared previously or, when a variable is declared, the analysis ensure that the identifier is not in use in the current scope. In contrast, the type checking ensures that the type of a variable is consistent with the use of the variable.

describe in 4.2.2

Both analyses are implemented as separate **Listeners**. This is done to separate both classes and give them concrete purposes while making the overall structure of the compiler more modular. In turn, since both analyses are implemented separately, they currently each require a transversal of the parse tree making the semantic analysis more inefficient. However, this can easily be addressed by writing a **Listener** that

wraps both analyses, transverses the parse tree only once, and calls the corresponding function for each wrapped analysis.

4.3.1 Declaration Analysis

The declaration analysis reports any errors that caused by the use of identifiers in invalid contexts; these are the use of an undeclared variable, the declaration of a variable in a context where it is already defined, and the use of a qubit in a code block which is guarded by itself. Additionally, it also reports a warning when a variable is declared but never used.

Define when a code block is guarded by a qubit.

For its analysis, the class creates a symbol table and adds any symbol that is declared. Further, it executes all other functions required for a valid analysis with the symbol table, like pushing scopes onto and popping them from the stack in the symbol table at the locations in the parse tree transversal where it is required. Before adding a symbol to the table, i.e. every time the parse tree transversal exits a declaration, the `Listener` checks whether the identifier is already declared in the current context; if this is the case, the corresponding error is reported. Similarly, each time an identifier is referenced outside of a declaration, the analysis ensures that the variable is defined, otherwise it reports the corresponding error as well. Lastly, whenever a register rule is encountered in the transversal, the analysis not only checks that the variable is declared but also whether the current code block is guarded by the referenced register. If this is the case, the analysis reports an error. While the register parsing rule is only used for gate applications and if-statements and not for function calls such as `sizeof`, only performing the check the register rule cannot result in an invalid use of an identifier. This is the case because all function calls result in constant values and only refer to properties of a register, not the register itself; therefore, the compiled program will not have a reference to the register in the controlled context.

Cannot check for all cases, some need to be at generation time

While the symbol table suffices for reporting the previous errors, it does not track the usage of symbols. Therefore, it cannot be used for the unused symbol warning. For this analysis, the class contains a dictionary that maps a symbol to the number of references. Each time a symbol is added to the symbol table, the corresponding entry in the dictionary is initialized to zero. Then, for every reference to the symbol, its usage counter is incremented. At the end of the tree transversal, specifically when exiting the main code block, the analysis iterates over all entries in the dictionary and reports a warning whenever the usage counter is still 0. However, there is an exception if an identifier is the underscore character so that it can be used as a throwaway variable, similar to other programming language.

4.3.2 Type Checking

... type checking and different errors it reports

4.3.3 Error Handling

- Describe the relevant parts of error handling (used by the semantic analysis)

4 Implementation

- Error handlers,
- How are errors reported ...
- error context

4.4 Code Generation

- How is code generated?
- Important classes and abstractions

4.4.1 Expressions

4.4.2 Composite Gates

4.5 Optimization

The implementation of the compiler does not only translate the custom language to OpenQASM 3 but and allows for optimization on the translated circuit. To apply the optimization to the translated circuit, the circuit description, i.e. the program, is used to build a circuit graph, as described in Sec. 3.3.2. Next, an algorithm iterates over the graph and checks whether a list of optimization rules is applicable to a part of the graph. If a rule is applicable, the rule is applied. The process of iterating over the entire graph is repeated for as long as rules were applied in the previous iteration over the graph. When the optimization of the circuit is completed, the graph is translated back to a programmatic description of the circuit and the result is returned.

In the following, we will discuss the implementation of the different steps in the optimization process. This includes the circuit graph in general, the construction of the graph based on the program, and the translation of the graph back to a circuit. Further, we discuss the implementation of the optimization rules and the optimization algorithm in general.

4.5.1 Circuit Graph

- Basic structure of circuit
- Graph construction
- Graph translation
- Some auxillary constructs and functions
 - Paths
 - Removal of nodes
 - replacing paths

Graph construction

Graph translation

- Important attributes
 - Gates applied in correct order
 - But when to apply which gate?
 - In many cases arbitrary (example)
- Eager translation
 - Translate each wire as much as possible
 - Switch to other wire only if entirely translated
 - or node can only be translated if other wire is translated first (up to the node)

4.5.2 Optimization Rules

- Rule interface
- Abstract optimization rule
- Describe the different optimization and the general implementation

4.5.3 Optimization Algorithm

- How is the graph iterated
- How are sub-paths used and created

4.6 Testing and Continuous Integration

- Different test categories
- How are they implemented?
- What do they test?
- (Continuous integration)

5 Conclusion and Future Work

- Conclusion to thesis
- Future work
 - how could language be extended

Bibliography

- [Aaby03] Anthony A. Aaby. *Compiler construction using flex and bison*. 2003.
- [AlGr05] T. Altenkirch and J. Grattage. A functional quantum programming language. In *20th Annual IEEE Symposium on Logic in Computer Science (LICS' 05)*, pages 249–258. IEEE, 2005. ISBN 0-7695-2266-1. doi: 10.1109/LICS.2005.1.
- [AGY07] Holger Bock Axelsen, Robert Glück, and Tetsuo Yokoyama. Reversible machine code and its abstract processor architecture. In Volker Diekert, Mikhail V. Volkov, and Andrei Voronkov, editors, *Computer Science – Theory and Applications*, volume 4649 of *Lecture Notes in Computer Science*, pages 56–69. Springer Berlin Heidelberg, Berlin, Heidelberg, 2007. ISBN 978-3-540-74509-9. doi: 10.1007/978-3-540-74510-5_9.
- [BGB*18] Ryan Babbush, Craig Gidney, Dominic W. Berry, Nathan Wiebe, Jarrod McClean, Alexandru Paler, Austin Fowler, and Hartmut Neven. Encoding electronic spectra in quantum circuits with linear t complexity. *Physical Review X*, 8(4), 2018. doi: 10.1103/PhysRevX.8.041015.
- [BFA22] Medina Bandic, Sebastian Feld, and Carmen G. Almudever. Full-stack quantum computing systems in the nisc era: algorithm-driven and hardware-aware compilation techniques. In Cristiana Bolchini, editor, *Proceedings of the 2022 Conference et Exhibition on Design, Automation et Test in Europe*, ACM Conferences, pages 1–6. European Design and Automation Association, Leuven,Belgium, 2022. ISBN 978-3-9819263-6-1. doi: 10.23919/DATE54114.2022.9774643.
- [BeLa17] Daniel J. Bernstein and Tanja Lange. Post-quantum cryptography. *Nature*, 549(7671):188–194, 2017. doi: 10.1038/nature23461.
- [BeVa93] Ethan Bernstein and Umesh Vazirani. Quantum complexity theory. In Rao Kosaraju, David Johnson, and Alok Aggarwal, editors, *Proceedings of the twenty-fifth annual ACM symposium on Theory of computing - STOC '93*, pages 11–20. ACM Press, New York, New York, USA, 1993. ISBN 0897915917. doi: 10.1145/167088.167097.
- [BBGV20] Benjamin Bichsel, Maximilian Baader, Timon Gehr, and Martin Vechev. Silq: a high-level quantum language with safe uncomputation and intuitive semantics. In Alastair F. Donaldson and Emina Torlak, editors, *Proceedings of the 41st ACM SIGPLAN Conference on Programming Language*

- Design and Implementation*, pages 286–300. ACM, New York, NY, USA, 2020. ISBN 9781450376136. doi: 10.1145/3385412.3386007.
- [BrBr02] Jean-Luc Brylinski and Raneer Brylinski. Universal quantum gates. In Goong Chen and Raneer Brylinski, editors, *Mathematics of Quantum Computation*, volume 20022356 of *Computational Mathematics*. Chapman and Hall/CRC, 2002. ISBN 978-1-58488-282-4. doi: 10.1201/9781420035377.pt2.
- [Copp02] D. Coppersmith. An approximate fourier transform useful in quantum factoring, 2002. doi: 10.48550/arXiv.quant-ph/0201067.
- [CJA*22] Andrew Cross, Ali Javadi-Abhari, Thomas Alexander, Niel de Beaudrap, Lev S. Bishop, Steven Heidel, Colm A. Ryan, Prasahnt Sivarajah, John Smolin, Jay M. Gambetta, and Blake R. Johnson. Openqasm 3: A broader and deeper quantum assembly language. *ACM Transactions on Quantum Computing*, 3(3):1–50, 2022. doi: 10.1145/3505636.
- [CBSG17] Andrew W. Cross, Lev S. Bishop, John A. Smolin, and Jay M. Gambetta. Open quantum assembly language, 2017. doi: 10.48550/arXiv.1707.03429.
- [Oliv07] José de Oliveira Guimarães. Learning compiler construction by examples. *ACM SIGCSE Bulletin*, 39(4):70–74, 2007. doi: 10.1145/1345375.1345418.
- [DeRe74] Franklin L. DeRemer. Lexical analysis. In F. L. Bauer, F. L. de Remer, M. Griffiths, U. Hill, J. J. Horning, C. H. A. Koster, W. M. McKeeman, P. C. Poole, W. M. Waite, and J. Eickel, editors, *Compiler Construction*, volume 21 of *Lecture Notes in Computer Science*, pages 109–120. Springer Berlin Heidelberg, Berlin, Heidelberg, 1974. ISBN 978-3-540-06958-4. doi: 10.1007/978-3-662-21549-4_5.
- [DeJo92] David Deutsch and Richard Jozsa. Rapid solution of problems by quantum computation. *Proceedings of the Royal Society of London. Series A: Mathematical and Physical Sciences*, 439(1907):553–558, 1992. doi: 10.1098/rspa.1992.0167.
- [DMN13] Simon J. Devitt, William J. Munro, and Kae Nemoto. Quantum error correction for beginners. *Reports on progress in physics. Physical Society (Great Britain)*, 76(7):076001, 2013. doi: 10.1088/0034-4885/76/7/076001.
- [Dijk75] Edsger W. Dijkstra. Guarded commands, nondeterminacy and formal derivation of programs. *Communications of the ACM*, 18(8):453–457, 1975. doi: 10.1145/360933.360975.
- [DiCh20b] Yongshan Ding and Frederic T. Chong. Introduction. In Yongshan Ding and Frederic T. Chong, editors, *Quantum Computer Systems*, Synthesis Lectures on Computer Architecture, pages 3–12.

- Springer International Publishing, Cham, 2020. ISBN 978-3-031-00637-1. doi: 10.1007/978-3-031-01765-0_1.
- [DiCh20c] Yongshan Ding and Frederic T. Chong. Quantum application design. In Yongshan Ding and Frederic T. Chong, editors, *Quantum Computer Systems*, Synthesis Lectures on Computer Architecture, pages 55–70. Springer International Publishing, Cham, 2020. ISBN 978-3-031-00637-1. doi: 10.1007/978-3-031-01765-0_3.
- [DiCh20a] Yongshan Ding and Frederic T. Chong. Think quantumly about computing. In Yongshan Ding and Frederic T. Chong, editors, *Quantum Computer Systems*, Synthesis Lectures on Computer Architecture, pages 13–54. Springer International Publishing, Cham, 2020. ISBN 978-3-031-00637-1. doi: 10.1007/978-3-031-01765-0_2.
- [DoSt99] Charles Donnelly and Richard Stallman. *Bison manual: The YACC-compatible parser generator, 3 November 1999, Bison Version 1.29*. Free Software Foundation, Boston, Mass., 1999. ISBN 1-882114-44-2.
- [Drap00] Thomas G. Draper. Addition on a quantum computer, 2000. doi: 10.48550/arXiv.quant-ph/0008033.
- [FNML21] Thomas Fösel, Murphy Yuezhen Niu, Florian Marquardt, and Li Li. Quantum circuit optimization with deep reinforcement learning, 2021. doi: 10.48550/arXiv.2103.07585.
- [GFH82] Mahadevan Ganapathi, Charles N. Fischer, and John L. Hennessy. Retargetable compiler code generation. *ACM Computing Surveys*, 14(4):573–592, 1982. doi: 10.1145/356893.356897.
- [GaCh11] Juan Carlos Garcia-Escartin and Pedro Chamorro-Posada. Equivalent quantum circuits, 2011. doi: 10.48550/arXiv.1110.2998.
- [Gref99] Gregory Grefenstette. Tokenization. In Nancy Ide, Jean Véronis, and Hans van Halteren, editors, *Syntactic Wordclass Tagging*, volume 9 of *Text, Speech and Language Technology*, pages 117–133. Springer Netherlands, Dordrecht, 1999. ISBN 978-90-481-5296-4. doi: 10.1007/978-94-015-9273-4_9.
- [HHHH09] Ryszard Horodecki, Paweł Horodecki, Michał Horodecki, and Karol Horodecki. Quantum entanglement. *Reviews of Modern Physics*, 81(2):865–942, 2009. doi: 10.1103/RevModPhys.81.865.
- [HuLe99] J. C. Huang and T. Leng. Generalized loop-unrolling: a method for program speedup. In *Proceedings 1999 IEEE Symposium on Application-Specific Systems and Software Engineering and Technology. ASSET’99 (Cat. No.PR00122)*, pages 244–248. IEEE Comput. Soc, 1999. ISBN 0-7695-0122-2. doi: 10.1109/ASSET.1999.756775.

Bibliography

- [Jozs05] Richard Jozsa. An introduction to measurement based quantum computation, 2005. doi: 10.48550/arXiv.quant-ph/0508124.
- [KMO*23] Fabian Kreppel, Christian Melzer, Diego Olvera Millán, Janis Wagner, Janine Hilder, Ulrich Poschinger, Ferdinand Schmidt-Kaler, and André Brinkmann. Quantum circuit compiler for a shuttling-based trapped-ion quantum computer. *Quantum*, 7:1176, 2023. doi: 10.22331/q-2023-11-08-1176.
- [KuBr00] Arun Kumar Pati and Samuel L. Braunstein. Impossibility of deleting an unknown quantum state. *Nature*, 404(6774):164–165, 2000. doi: 10.1038/404130b0.
- [Land61] R. Landauer. Irreversibility and heat generation in the computing process. *IBM Journal of Research and Development*, 5(3):183–191, 1961. doi: 10.1147/rd.53.0183.
- [LPM*24] Zikun Li, Jinjun Peng, Yixuan Mei, Sina Lin, Yi Wu, Oded Padon, and Zhihao Jia. Quarl: A learning-based quantum circuit optimizer. *Proceedings of the ACM on Programming Languages*, 8(OOPSLA1):555–582, 2024. doi: 10.1145/3649831.
- [LBZ21] Ji Liu, Luciano Bello, and Huiyang Zhou. Relaxed peephole optimization: A novel compiler optimization for quantum circuits. In *Proceedings of the 2021 IEEE/ACM International Symposium on Code Generation and Optimization*, pages 301–314. IEEE Press, [S.l.], 2021. ISBN 978-1-7281-8613-9. doi: 10.1109/CGO51591.2021.9370310.
- [LoCh19] Guang Hao Low and Isaac L. Chuang. Hamiltonian simulation by qubitization. *Quantum*, 3:163, 2019. doi: 10.22331/q-2019-07-12-163.
- [MVZJ18] Vasileios Mavroeidis, Kamer Vishi, Mateusz D. Zych, and Audun Jøsang. The impact of quantum computing on present cryptography. 2018. doi: 10.48550/arXiv.1804.00200.
- [McKe65] W. M. McKeeman. Peephole optimization. *Communications of the ACM*, 8(7):443–444, 1965. doi: 10.1145/364995.365000.
- [MHH19] Gary J. Mooney, Charles D. Hill, and Lloyd C. L. Hollenberg. Entanglement in a 20-qubit superconducting quantum computer. *Scientific reports*, 9(1):13465, 2019. doi: 10.1038/s41598-019-49805-7.
- [Niel06] Michael A. Nielsen. Cluster-state quantum computation. *Reports on Mathematical Physics*, 57(1):147–161, 2006. doi: 10.1016/S0034-4877(06)80014-5.
- [PaQu95] T. J. Parr and R. W. Quong. Antlr: A predicated-ll(k) parser generator. *Software: Practice and Experience*, 25(7):789–810, 1995. doi: 10.1002/spe.4380250705.

- [PaFi11] Terence Parr and Kathleen Fisher. Ll(*): the foundation of the antlr parser generator. *ACM SIGPLAN Notices*, 46(6):425–436, 2011. doi: 10.1145/1993316.1993548.
- [PeMa02] SIMON PEYTON JONES and SIMON MARLOW. Secrets of the glasgow haskell compiler inliner. *Journal of Functional Programming*, 12(4-5):393–434, 2002. doi: 10.1017/S0956796802004331.
- [Pres18] John Preskill. Quantum computing in the nisq era and beyond. *Quantum*, 2:79, 2018. doi: 10.22331/q-2018-08-06-79.
- [RDB*22] Roman Rietsche, Christian Dremel, Samuel Bosch, Léa Steinacker, Miriam Meckel, and Jan-Marco Leimeister. Quantum computing. *Electronic Markets*, 32(4):2525–2536, 2022. doi: 10.1007/s12525-022-00570-y.
- [RSA78] R. L. Rivest, A. Shamir, and L. Adleman. A method for obtaining digital signatures and public-key cryptosystems. *Communications of the ACM*, 21(2):120–126, 1978. doi: 10.1145/359340.359342.
- [RLB*24] Francisco J. R. Ruiz, Tuomas Laakkonen, Johannes Bausch, Matej Balog, Mohammadamin Barekatain, Francisco J. H. Heras, Alexander Novikov, Nathan Fitzpatrick, Bernardino Romera-Paredes, John van de Wetering, Alhussein Fawzi, Konstantinos Meichanetzidis, and Pushmeet Kohli. Quantum circuit optimization with alphasensor, 2024. doi: 10.48550/arXiv.2402.14396.
- [SWW*88] V. Seshadri, S. Weber, D. B. Wortman, C. P. Yu, and I. Small. Semantic analysis in a concurrent compiler. In R. L. Wexelblat, editor, *Proceedings of the ACM SIGPLAN 1988 conference on Programming language design and implementation*, pages 233–240. ACM, New York, NY, USA, 1988. ISBN 0897912691. doi: 10.1145/53990.54013.
- [Shor97] Peter W. Shor. Polynomial-time algorithms for prime factorization and discrete logarithms on a quantum computer. *SIAM Journal on Computing*, 26(5):1484–1509, 1997. doi: 10.1137/S0097539795293172.
- [TvS82] Andrew S. Tanenbaum, Hans van Staveren, and Johan W. Stevenson. Using peephole optimization on intermediate code. *ACM Transactions on Programming Languages and Systems*, 4(1):21–36, 1982. doi: 10.1145/357153.357155.
- [TGS22] Theodoros Theodoridis, Tobias Grosser, and Zhendong Su. Understanding and exploiting optimal function inlining. In Babak Falsafi, Michael Ferdman, Shan Lu, and Tom Wenisch, editors, *Proceedings of the 27th ACM International Conference on Architectural Support for Programming Languages and Operating Systems*, pages 977–989. ACM, New York, NY, USA, 2022. ISBN 9781450392051. doi: 10.1145/3503222.3507744.

- [TAG12] Michael Kirkedal Thomsen, Holger Bock Axelsen, and Robert Glück. A reversible processor architecture and its reversible logic design. In David Hutchison, Takeo Kanade, Josef Kittler, Jon M. Kleinberg, Friedemann Mattern, John C. Mitchell, Moni Naor, Oscar Nierstrasz, C. Pandu Rangan, Bernhard Steffen, Madhu Sudan, Demetri Terzopoulos, Doug Tygar, Moshe Y. Vardi, Gerhard Weikum, Alexis de Vos, and Robert Wille, editors, *Reversible Computation*, volume 7165 of *Lecture Notes in Computer Science*, pages 30–42. Springer Berlin Heidelberg, Berlin, Heidelberg, 2012. ISBN 978-3-642-29516-4. doi: 10.1007/978-3-642-29517-1_3.
- [VSSD07] Alfred V. Aho, Monica S. Lam, Ravi Sethi, and Jeffrey D. Ullmann. *Compilers: Principles, techniques, & tools*. Pearson/Addison Wesley, Boston, 2nd ed. edition, 2007. ISBN 0-321-48681-1.
- [Wait74] W. M. Waite. Semantic analysis. In F. L. Bauer, F. L. de Remer, M. Griffiths, U. Hill, J. J. Horning, C. H. A. Koster, W. M. McKeeman, P. C. Poole, W. M. Waite, and J. Eickel, editors, *Compiler Construction*, volume 21 of *Lecture Notes in Computer Science*, pages 157–169. Springer Berlin Heidelberg, Berlin, Heidelberg, 1974. ISBN 978-3-540-06958-4. doi: 10.1007/978-3-662-21549-4_8.
- [WeZa91] Mark N. Wegman and F. Kenneth Zadeck. Constant propagation with conditional branches. *ACM Transactions on Programming Languages and Systems*, 13(2):181–210, 1991. doi: 10.1145/103135.103136.
- [Wino78] S. Winograd. On computing the discrete fourier transform. *Mathematics of Computation*, 32(141):175–199, 1978. doi: 10.1090/S0025-5718-1978-0468306-4.
- [WoZu82] W. K. Wootters and W. H. Zurek. A single quantum cannot be cloned. *Nature*, 299(5886):802–803, 1982. doi: 10.1038/299802a0.
- [Ying11] Mingsheng Ying. Floyd–hoare logic for quantum programs. *ACM Transactions on Programming Languages and Systems*, 33(6):1–49, 2011. doi: 10.1145/2049706.2049708.
- [YYF12] Mingsheng Ying, Nengkun Yu, and Yuan Feng. Defining quantum control flow, 2012. doi: 10.48550/arXiv.1209.4379.
- [YuCa22] Charles Yuan and Michael Carbin. Tower: data structures in quantum superposition. *Proceedings of the ACM on Programming Languages*, 6(OOPSLA2):259–288, 2022. doi: 10.1145/3563297.
- [YVC24] Charles Yuan, Agnes Villanyi, and Michael Carbin. Quantum control machine: The limits of control flow in quantum programming. *Proceedings of the ACM on Programming Languages*, 8(OOPSLA1):1–28, 2024. doi: 10.1145/3649811.

- [ZLY17] Ye Zhang, Yuliang Lu, and Bin Yang. Parsing statement list program using flex and bison. In *2017 First International Conference on Electronics Instrumentation & Information Systems (EIIS)*, pages 1–4. IEEE, 2017. ISBN 978-1-5386-0843-2. doi: 10.1109/EIIS.2017.8298547.

List of Figures

| | | |
|------|--|----|
| 2.1 | Null gates of self-inverse gates. | 7 |
| 2.2 | Null gates for gate in specific conditions. | 8 |
| 2.3 | Control reversal of the controlled Z gate. | 8 |
| 2.4 | Control reversal of CX. | 8 |
| 2.5 | A non-reversible exponentiation algorithm. | 13 |
| 2.6 | Reversible exponentiation algorithm. | 13 |
| 2.7 | A synchronized, reversible exponentiation algorithm. | 14 |
| 2.8 | Code for an OpenQASM 3 example circuit. | 16 |
| 2.9 | An example of a regular grammar for the lexical analysis. | 17 |
| 2.10 | An example of a context-free grammar for parsing simple expressions. | 18 |
| 2.11 | Simple ANTLR4 grammar for expressions. | 22 |
| 3.1 | Luie gate definition for the Quantum Fourier Transform. | 28 |
| 3.2 | An example of a simple, unoptimized circuit graph. | 34 |
| 3.3 | Circuit graph after the first optimization. | 35 |
| 3.4 | Completely optimized graph. | 35 |
| 3.5 | A command line interface example. | 37 |
| 4.1 | The basic structure parsing rules of Luie. | 39 |
| 4.2 | The parsing rules for declarations in Luie. | 40 |